



Grymoire Navigation

[Unix/Linux](#)
[Quotes](#)
[Bourne Shell](#)
[C Shell](#)
[File Permissions](#)
[Regular](#)
[Expressions](#)
[grep](#)
[awk](#) UPDATED
[sed](#) UPDATED
[find](#)
[tar](#)
[inodes](#)
[Security](#)
[IPv6](#)
[Wireless](#)
[Hardware](#)
[spam](#)
[Deception](#)
[PostScript](#)
[Halftones](#)
[Privacy](#)
[Bill of Rights](#)
[References](#)
[Top 10 reasons to avoid CSH](#)
[sed Chart](#) PDF
[awk Reference](#)
[HTML](#)
[Magic](#)
[Search](#)
[About](#)
[Donate](#) NEW

Google+: [Bruce Barnett](#)

Twitter: [@grymoire](#)

Blog: [Wordpress](#)

Sed - An Introduction and Tutorial

by Bruce Barnett

Last modified: Thu Apr 23 16:37:48 EDT 2015

Quick Links

Sed Commands		
: label	# comment	{....} Block
= - print line number	a \ - Append	b label - Branch
c \ - change	d and D - Delete	g and G - Get
h and H - Hold	i \ - Insert	l - Look
n and N - Next	p and P - Print	q - Quit
r filename - Read File	s/.../.../ - Substitute	t label - Test
w filename - Write Filename	x - eXchange	y/.../.../ - Transform

Sed Pattern Flags
/g - Global
/I - Ignore Case
/p - Print
/w filename - Write Filename

Sed Command Line options	
Short Option (Long option)	Sed version
-n	Classic
-e script	Classic
-f scriptfile	Classic
-e script (--expression=script)	GNU sed
-f scriptfile (--file=scriptfile)	GNU sed
-h (--help)	GNU sed
-n (--quiet --silent)	GNU sed
-V (--version)	GNU sed
-r (--regexp-extended)	GNU sed
-i[SUFFIX] (--in-place[=SUFFIX])	GNU sed
-l N (--line-length=N)	GNU sed
-b (--binary)	GNU sed
-s (--separate)	GNU sed
-z (--null-data)	GNU sed
-u (--unbuffered)	GNU sed
(--follow-symlinks)	GNU sed
(--posix)	GNU sed
-i SUFFIX	Mac OS X, FreeBSD
-a	Mac OS X, FreeBSD
-l	Mac OS X, FreeBSD
-E	Mac OS X, FreeBSD
-r	FreeBSD
-I SUFFIX	FreeBSD

Table of Contents

Note - You can click on the table of contents sections to jump to that section.

Then click on the section header of any section to jump back to the table of contents.

- The Awful Truth about sed
- The essential command: s for substitution
- The slash as a delimiter
- Using & as the matched string
- Using \1 to keep part of the pattern
- Extended Regular Expressions
- Sed Pattern Flags
 - /g - Global replacement
- Is sed recursive?
 - /1, /2, etc. Specifying which occurrence
- /p - print
- Write to a file with /w filename
- /I - Ignore Case
- Combining substitution flags
- Arguments and invocation of sed
 - Multiple commands with -e command
 - Filenames on the command line
 - sed -n: no printing
 - Using 'sed /pattern/'
 - Using 'sed -n /pattern/p' to duplicate the function of grep
 - sed -f scriptname
 - sed in shell scripts
 - Quoting multiple sed lines in the C shell
 - Quoting multiple sed lines in the Bourne shell
 - sed -V
 - sed -h
 - A sed interpreter script
 - Sed Comments
 - Passing arguments into a sed script
 - Using sed in a shell here-is document
 - Multiple commands and order of execution
- Addresses and Ranges of Text
 - Restricting to a line number
 - Patterns
 - Ranges by line number
 - Ranges by patterns
- Delete with d
- Printing with p
- Reversing the restriction with !
- Relationships between d, p, and !
- The q or quit command
- Grouping with { and }
- Operating in a pattern range except for the patterns
- Writing a file with the 'w' command
- Reading in a file with the 'r' command
- SunOS and the # Comment Command
- Adding, Changing, Inserting new lines
 - Append a line with 'a'
 - Insert a line with 'i'
 - Change a line with 'c'
 - Leading tabs and spaces in a sed script
 - Adding more than one line
 - Adding lines and the pattern space
 - Address ranges and the above commands
- Multi-Line Patterns
- Print line number with =
- Transform with y
- Displaying control characters with a l
- Working with Multiple Lines
 - Matching three lines with sed
 - Matching patterns that span multiple lines
 - Using newlines in sed scripts
 - The Hold Buffer
 - Exchange with x
 - Example of Context Grep
 - Hold with h or H

Keeping more than one line in the hold buffer
Get with g or G
Flow Control
Testing with t
Debugging with l
An alternate way of adding comments
The poorly documented ;
Passing regular expressions as arguments
Inserting binary characters
GNU sed Command Line arguments
 The -posix argument
 The --version argument
 The -h Help argument
 The -l Line Length Argument
 The -s Separate argument
 The -i in-place argument
 The --follow-symlinks argument
 The -b Binary argument
 The -r Extended Regular Expression argument
 The -u Unbuffered argument
 The -z Null Data argument
FreeBSD Extensions
 -a or delayed open
 The -I in-place argument
 -E or Extended Regular Expressions
Using word boundaries
Command Summary
In Conclusion
More References

Copyright 1994, 1995 Bruce Barnett and General Electric Company

Copyright 2001,2005,2007,2011,2013
Bruce Barnett

All rights reserved

You are allowed to print copies of this tutorial for your personal use, and link to this page, but you are not allowed to make electronic copies, or redistribute this tutorial in any form without permission.

Original version written in 1994 and published in the Sun Observer

Introduction to Sed

How to use sed, a special editor for modifying files automatically. If you want to write a program to make changes in a file, sed is the tool to use.

There are a few programs that are the real workhorse in the UNIX toolbox. These programs are simple to use for simple applications, yet have a rich set of commands for performing complex actions. Don't let the complex potential of a program keep you from making use of the simpler aspects. I'll start with the simple concepts and introduce the advanced topics later on.

When I first wrote this (in 1994), most versions of `sed` did not allow you to place comments inside the script. Lines starting with the '#' characters are comments. Newer versions of `sed` may support comments at the end of the line as well.

One way to think of this is that the old, "classic" version was the basis of GNU, FreeBSD and Solaris versions of `sed`. And to help you understand what I had to work with, here is the [**sed\(1\) manual page from Sun/Oracle**](#) .

The Awful Truth about sed

Sed is the ultimate **s**tream **e**ditor. If that sounds strange, picture a stream flowing through a pipe. Okay, you can't see a stream if it's inside a pipe. That's what I get for attempting a flowing analogy. You want literature, read James Joyce.

Anyhow, *sed* is a marvelous utility. Unfortunately, most people never learn its real power. The language is very simple, but the documentation is terrible. The Solaris on-line manual pages for *sed* are five pages long, and two of those pages describe the 34 different errors you can get. A program that spends as much space documenting the errors as it does documenting the language has a

serious
learning
curve.
**Do
not
fret!**
It
is
not
your
fault
you
don't
understand
sed.
I
will
cover
sed
completely.
But
I
will
describe
the
features
in
the
order
that
I
learned
them.
I
didn't
learn
everything
at
once.
You
don't
need
to
either.

**The
essential
command:
s
for
substitution**

Sed
has
several
commands,
but
most
people
only
learn
the
substitute

command:

s.
The
substitute
command
changes
all
occurrences
of
the
regular
expression
into
a
new
value.
A
simple
example
is
changing
"day"
in
the
"old"
file
to
"night"
in
the
"new"
file:

```
sed s/day/night/ <old >new
```

Or
another
way
(for
UNIX
beginners),

```
sed s/day/night/ old >new
```

and
for
those
who
want
to
test
this:

```
echo day | sed s/day/night/
```

This
will
output
"night".

I
didn't
put
quotes
around
the
argument
because
this
example

```
didn't
need
them.
If
you
read
my
earlier
tutorial
on
quotes ,
you
would
understand
why
it
doesn't
need
quotes.
However,
I
recommend
you
do
use
quotes.
If
you
have
meta-characters
in
the
command,
quotes
are
necessary.
And
if
you
aren't
sure,
it's
a
good
habit,
and
I
will
henceforth
quote
future
examples
to
emphasize
the
"best
practice."
Using
the
strong
(single
quote)
character,
that
would
be:

sed 's/day/night/' <old >new
```

```
I
must
emphasize
that
the
sed
editor
changes
exactly
what
you
tell
it
to.
So
if
you
executed

echo Sunday | sed 's/day/night/'

This
would
output
the
word
"Sunnight"
because
sed
found
the
string
"day"
in
the
input.

Another
important
concept
is
that
sed
is
line
oriented.
Suppose
you
have
the
input
file:

one two three, one two three
four three two one
one hundred

and
you
used
the
command

sed 's/one/ONE/' <file

The
output
would
be

ONE two three, one two three
```



```
four three two ONE
ONE hundred

Note
that
this
changed
"one"
to
"ONE"
once
on
each
line.
The
first
line
had
"one"
twice,
but
only
the
first
occurrence
was
changed.
That
is
the
default
behavior.
If
you
want
something
different,
you
will
have
to
use
some
of
the
options
that
are
available.
I'll
explain
them
later.

So
let's
continue.

There
are
four
parts
to
this
substitute
command:

s          Substitute command
/.../     Delimiter
one       Regular Expression Pattern Search Pattern
```

ONE Replacement string

The
search
pattern
is
on
the
left
hand
side
and
the
replacement
string
is
on
the
right
hand
side.

We've
covered
quoting
and
regular
expressions. .

That's
90%
of
the
effort
needed
to
learn
the
substitute
command.

To
put
it
another
way,
you
already
know
how
to
handle
90%
of
the
most
frequent
uses
of
sed.
There
are
a
...
few
fine
points
that
any
future
sed

expert
should
know
about.
(You
just
finished
section
1.
There
are
only
63
more
sections
to
cover.
:-)
Oh.
And
you
may
want
to
bookmark
this
page,
....
just
in
case
you
don't
finish.

The slash as a delimiter

The
character
after
the
s
is
the
delimiter.
It
is
conventionally
a
slash,
because
this
is
what
ed,
more,
and
vi
use.
It
can

```
be
anything
you
want,
however.
If
you
want
to
change
a
pathname
that
contains
a
slash
-
say
/usr/local/bin
to
/common/bin
-
you
could
use
the
backslash
to
quote
the
slash:

sed 's/\\usr\\local\\bin\\/common\\bin/' <old >new

Gulp.
Some
call
this
a
'Picket
Fence'
and
it's
ugly.
It
is
easier
to
read
if
you
use
an
underline
instead
of
a
slash
as
a
delimiter:

sed 's/_usr/local/bin_/common/bin_' <old >new

Some
people
use
colons:

sed 's:/usr/local/bin:/common/bin:' <old >new
```

Others
use
the
"`|`"
character.

```
sed 's|/usr/local/bin|/common/bin|' <old >new
```

Pick
one
you
like.
As
long
as
it's
not
in
the
string
you
are
looking
for,
anything
goes.
And
remember
that
you
need
three
delimiters.
If
you
get
a
"Unterminated
's'
command"
it's
because
you
are
missing
one
of
them.

Using & as the matched string

Sometimes
you
want
to
search
for
a
pattern
and

```
add
some
characters,
like
parenthesis,
around
or
near
the
pattern
you
found.
It
is
easy
to
do
this
if
you
are
looking
for
a
particular
string:

sed 's/abc/(abc)/' <old >new
```

```
This
won't
work
if
you
don't
know
exactly
what
you
will
find.
How
can
you
put
the
string
you
found
in
the
replacement
string
if
you
don't
know
what
it
is?
```

```
The
solution
requires
the
special
character
"&."
It
```

corresponds
to
the
pattern
found.

```
sed 's/[a-z]*/(&)/' <old >new
```

You
can
have
any
number
of
"&"
in
the
replacement
string.
You
could
also
double
a
pattern,
e.g.
the
first
number
of
a
line:

```
% echo "123 abc" | sed 's/[0-9]*/& &/'  
123 123 abc
```

Let
me
slightly
amend
this
example.
Sed
will
match
the
first
string,
and
make
it
as
greedy
as
possible.
I'll
cover
that
later.
If
you
don't
want
it
to
be
so
greedy
(i.e.
limit

the
matching),
you
need
to
put
restrictions
on
the
match.

The
first
match
for
'[0-9]*'
is
the
first
character
on
the
line,
as
this
matches
zero
or
more
numbers.

So
if
the
input
was
"abc
123"
the
output
would
be
unchanged
(well,
except
for
a
space
before
the
letters).

A
better
way
to
duplicate
the
number
is
to
make
sure
it
matches
a
number:

```
% echo "123 abc" | sed 's/[0-9][0-9]*/& &/'  
123 123 abc
```


The string "abc" is unchanged, because it was not matched by the regular expression. If you wanted to eliminate "abc" from the output, you must expand the regular expression to match the rest of the line and explicitly exclude part of the expression using "(" and "\1", which is the next topic.

Extended Regular Expressions

Let me add a quick comment here because

there
is
another
way
to
write
the
above
script.
"[0-9]*"
matches
zero
or
more
numbers.
"[0-9]
[0-9]*"
matches
one
or
more
numbers.
Another
way
to
do
this
is
to
use
the
"+"
meta-character
and
use
the
pattern
"[0-9]+"as
the
"+"
is
a
special
character
when
using
"extended
regular
expressions."
Extended
regular
expressions
have
more
power,
but
sed
scripts
that
treated
"+"
as
a
normal
character
would
break.
Therefore

you
must
explicitly
enable
this
extension
with
a
command
line
option.

GNU
sed
turns
this
feature
on
if
you
use
the
"-r"
command
line
option.
So
the
above
could
also
be
written
using

```
% echo "123 abc" | sed -r 's/[0-9]+/& &/'  
123 123 abc
```

Mac
OS
X
and
FreeBSD
uses

-E
instead
of

-r .

For
more
information
on
extended
regular
expressions,
see

**Regular
Expressions**

and
the
description

of
the

-r
command
line
argument

Using

\1 to keep part of the pattern

I
have
already
described
the
use
of
"("
")"
and
"1"
in
my
tutorial
on
regular
expressions.
To
review,
the
escaped
parentheses
(that
is,
parentheses
with
backslashes
before
them)
remember
a
substring
of
the
characters
matched
by
the
regular
expression.
You
can
use
this
to
exclude
part
of
the
characters
matched
by
the
regular
expression.
The

```
"\1"
is
the
first
remembered
pattern,
and
the
"\2"
is
the
second
remembered
pattern.
Sed
has
up
to
nine
remembered
patterns.

If
you
wanted
to
keep
the
first
word
of
a
line,
and
delete
the
rest
of
the
line,
mark
the
important
part
with
the
parenthesis:

sed 's/\([a-z]*\) */\1/'

I
should
elaborate
on
this.
Regular
expressions
are
greedy,
and
try
to
match
as
much
as
possible.
"[a-z]*"
matches
```

zero
or
more
lower
case
letters,
and
tries
to
match
as
many
characters
as
possible.

The
".*"
matches
zero
or
more
characters
after
the
first
match.
Since
the
first
one
grabs
all
of
the
contiguous
lower
case
letters,
the
second
matches
anything
else.
Therefore
if
you
type

```
echo abcd123 | sed 's/\([a-z]*\)*/\1/'
```

This
will
output
"abcd"
and
delete
the
numbers.

If
you
want
to
switch
two
words
around,
you
can

remember
two
patterns
and
change
the
order
around:

```
sed 's/\([a-z]*\) \([a-z]*\)/\2 \1/'
```

Note
the
space
between
the
two
remembered
patterns.
This
is
used
to
make
sure
two
words
are
found.
However,
this
will
do
nothing
if
a
single
word
is
found,
or
any
lines
with
no
letters.
You
may
want
to
insist
that
words
have
at
least
one
letter
by
using

```
sed 's/\([a-z][a-z]*\) \([a-z][a-z]*\)/\2 \1/'
```

or
by
using
extended
regular
expressions
(note

```
that
'('
and
')'
no
longer
need
to
have
a
backslash):

sed -r 's/([a-z]+) ([a-z]+)/\2 \1/' # Using GNU sed
sed -E 's/([a-z]+) ([a-z]+)/\2 \1/' # Using Apple Mac OS X
```

```
The
"\1"
doesn't
have
to
be
in
the
replacement
string
(in
the
right
hand
side).
It
can
be
in
the
pattern
you
are
searching
for
(in
the
left
hand
side).
If
you
want
to
eliminate
duplicated
words,
you
can
try:

sed 's/\([a-z]*\) \1/\1/'

If
you
want
to
detect
duplicated
words,
you
can
use

sed -n '/\([a-z][a-z]*\) \1/p'
```


or
with
extended
regular
expressions

```
sed -rn '/([a-z]+) \1/p' # GNU sed  
sed -En '/([a-z]+) \1/p' # Mac OS X
```

This,
when
used
as
a
filter,
will
print
lines
with
duplicated
words.

The
numeric
value
can
have
up
to
nine
values:
"1"
thru
"9."
If
you
wanted
to
reverse
the
first
three
characters
on
a
line,
you
can
use

```
sed 's/^(.)\)(.)\)(.)\)/\3\2\1/'
```

Sed Pattern Flags

You
can
add
additional
flags
after
the
last
delimiter.
You
might

have noticed I used a 'p' at the end of the previous substitute command. I also added the '-n' option. Let me first cover the 'p' and other pattern flags. These flags can specify what happens when a match is found. Let me describe them.

/g
=
Global replacement

Most UNIX utilities work on files, reading a line at a time. *Sed*, by default,

is
the
same
way.
If
you
tell
it
to
change
a
word,
it
will
only
change
the
first
occurrence
of
the
word
on
a
line.
You
may
want
to
make
the
change
on
every
word
on
the
line
instead
of
the
first.
For
an
example,
let's
place
parentheses
around
words
on
a
line.
Instead
of
using
a
pattern
like
"[A-Za-z]*"
which
won't
match
words
like
"won't,"
we
will
use

a
pattern,
"[
^
]*,"
that
matches
everything
except
a
space.
Well,
this
will
also
match
anything
because
"*"
means
**zero
or
more.**
The
current
version
of
Solaris's
sed
(as
I
wrote
this)
can
get
unhappy
with
patterns
like
this,
and
generate
errors
like
"Output
line
too
long"
or
even
run
forever.
I
consider
this
a
bug,
and
have
reported
this
to
Sun.
As
a
work-around,
you
must
avoid
matching

the
null
string
when
using
the
"g"
flag
to
sed.
A
work-around
example
is:
"[^
][[^
]*."
The
following
will
put
parenthesis
around
the
first
word:

sed 's/[^]*/(&)/' <old >new

If
you
want
it
to
make
changes
for
every
word,
add
a
"g"
after
the
last
delimiter
and
use
the
work-around:

sed 's/[^][^]*/(&)/g' <old >new

Is sed recursive?

Sed
only
operates
on
patterns
found
in
the
in-coming
data.

That
is,
the
input
line
is
read,
and
when
a
pattern
is
matched,
the
modified
output
is
generated,
and
the
rest
of
the
input
line
is
scanned.
The
"s"
command
will
not
scan
the
newly
created
output.
That
is,
you
don't
have
to
worry
about
expressions
like:

```
sed 's/loop/loop the loop/g' <old >new
```

This
will
not
cause
an
infinite
loop.
If
a
second
"s"
command
is
executed,
it
could
modify
the
results

of
a
previous
command.
I
will
show
you
how
to
execute
multiple
commands
later.

/1,
/2,
etc.
Specifying
which
occurrence

With
no
flags,
the
first
matched
substitution
is
changed.
With
the
"g"
option,
all
matches
are
changed.
If
you
want
to
modify
a
particular
pattern
that
is
not
the
first
one
on
the
line,
you
could
use
"\
and
"\
to
mark

each
pattern,
and
use
"\1"
to
put
the
first
pattern
back
unchanged.
This
next
example
keeps
the
first
word
on
the
line
but
deletes
the
second:

```
sed 's/\([a-zA-Z]*\) \([a-zA-Z]*\) /\1 /' <old >new
```

Yuck.
There
is
an
easier
way
to
do
this.
You
can
add
a
number
after
the
substitution
command
to
indicate
you
only
want
to
match
that
particular
pattern.
Example:

```
sed 's/[a-zA-Z]* //2' <old >new
```

You
can
combine
a
number
with
the
g
(global)

flag.
For
instance,
if
you
want
to
leave
the
first
word
alone,
but
change
the
second,
third,
etc.
to
be
DELETED
instead,
use
/2g:

```
sed 's/[a-zA-Z]* /DELETED /2g' <old >new
```

I've
heard
that
combining
the
number
with
the
g
command
does
not
work
on
The
MacOS,
and
perhaps
the
FreeSBD
version
of
sed
as
well.

Don't
get
/2
and
\2
confused.
The
/2
is
used
at
the
end.
\2
is
used

in
inside
the
replacement
field.

Note
the
space
after
the
"*" character.

Without
the
space,
sed
will
run
a
long,
long
time.

(Note:
this
bug
is
probably
fixed
by
now.)

This
is
because
the
number
flag
and
the
"g"
flag
have
the
same
bug.
You
should
also
be
able
to
use
the
pattern

```
sed 's/[^ ]*/2' <old >new
```

but
this
also
eats
CPU.
If
this
works
on
your
computer,
and

it
does
on
some
UNIX
systems,
you
could
remove
the
encrypted
password
from
the
password
file:

```
sed 's/[^:]*//2' </etc/passwd >/etc/password.new
```

But
this
didn't
work
for
me
the
time
I
wrote
this.
Using
"`[^:][^:]*`"
as
a
work-around
doesn't
help
because
it
won't
match
an
non-existent
password,
and
instead
delete
the
third
field,
which
is
the
user
ID!
Instead
you
have
to
use
the
ugly
parenthesis:

```
sed 's/^\([^:]*\):[^:]:/\1:/' </etc/passwd >/etc/password.new
```

You
could
also
add

a
character
to
the
first
pattern
so
that
it
no
longer
matches
the
null
pattern:

```
sed 's/[^:]*:/:/2' </etc/passwd >/etc/password.new
```

The
number
flag
is
not
restricted
to
a
single
digit.
It
can
be
any
number
from
1
to
512.
If
you
wanted
to
add
a
colon
after
the
80th
character
in
each
line,
you
could
type:

```
sed 's/./&:/80' <file >new
```

You
can
also
do
it
the
hard
way
by
using
80
dots:

```
sed 's/^...../&:/' <
```

/p
=
print

By default, *sed* prints every line. If it makes a substitution, the new text is printed instead of the old one. If you use an optional argument to *sed*, "sed -n," it will not, by default, print any new lines. I'll cover this and other options later. When the "-n" option is used, the "p" flag will cause the modified

line
to
be
printed.
Here
is
one
way
to
duplicate
the
function
of
grep
with
sed:

```
sed -n 's/pattern/&/p' <file
```

But
a
simpler
version
is
described

later

Write
to
a
file
with
/w
filename

There
is
one
more
flag
that
can
follow
the
third
delimiter.
With
it,
you
can
specify
a
file
that
will
receive
the
modified
data.
An
example
is
the
following,

which
will
write
all
lines
that
start
with
an
even
number,
followed
by
a
space,
to
the
file
even:

```
sed -n 's/^[0-9]*[02468] /&/w even' <file
```

In
this
example,
the
output
file
isn't
needed,
as
the
input
was
not
modified.
You
must
have
exactly
one
space
between
the
w
and
the
filename.
You
can
also
have
ten
files
open
with
one
instance
of
sed.
This
allows
you
to
split
up
a
stream
of

data
into
separate
files.
Using
the
previous
example
combined
with
multiple
substitution
commands
described
later,
you
could
split
a
file
into
ten
pieces
depending
on
the
last
digit
of
the
first
number.
You
could
also
use
this
method
to
log
error
or
debugging
information
to
a
special
file.

/I

=
Ignore
Case

GNU
has
added
another
pattern
flags

-
/I

This
flag
makes

the
pattern
match
case
insensitive.
This
will
match
abc,
aBc,
ABC,
AbC,
etc.:

```
sed -n '/abc/I p' <old >new
```

Note
that
a
space
after
the
'/I'
and
the
'p'
(print)
command
emphasizes
that
the
'p'
is
not
a
modifier
of
the
pattern
matching
process,
,
but
a
command
to
execute
after
the
pattern
matching.

Combining substitution flags

You
can
combine
flags
when
it
makes
sense.
Please
note
that

the
"w"
has
to
be
the
last
flag.
For
example
the
following
command
works:

```
sed -n 's/a/A/2pw /tmp/file' <old >new
```

Next
I
will
discuss
the
options
to
sed,
and
different
ways
to
invoke
sed.

Arguments and invocation of sed

previously,
I
have
only
used
one
substitute
command.
If
you
need
to
make
two
changes,
and
you
didn't
want
to
read
the
manual,
you
could
pipe
together

multiple
sed
commands:

`sed 's/BEGIN/begin/' <old | sed 's/END/end/' >new`

This
used
two
processes
instead
of
one.
A
sed
guru
never
uses
two
processes
when
one
can
do.

Multiple commands with -e command

One
method
of
combining
multiple
commands
is
to
use
a
-e
before
each
command:

`sed -e 's/a/A/' -e 's/b/B/' <old >new`

A
"-e"
isn't
needed
in
the
earlier
examples
because
sed
knows
that
there
must
always
be
one
command.

If
you
give
sed
one
argument,
it
must
be
a
command,
and
sed
will
edit
the
data
read
from
standard
input.

The
long
argument
version
is

```
sed --expression='s/a/A/' --expression='s/b/B/' <old >new
```

Also
see

Quoting
multiple
sed
lines
in
the
Bourne
shell

Filenames
on
the
command
line

You
can
specify
files
on
the
command
line
if
you
wish.
If
there
is
more
than
one
argument
to

sed
that
does
not
start
with
an
option,
it
must
be
a
filename.

This
next
example
will
count
the
number
of
lines
in
three
files
that
don't
begin
with
a
"#:"

```
sed 's/^#.*//' f1 f2 f3 | grep -v '^$' | wc -l
```

Let's
break
this
down
into
pieces.

The
sed
substitute
command
changes
every
line
that
starts
with
a
"#"

into
a
blank
line.

Grep
was
used
to
filter
out
(delete)
empty
lines.

Wc
counts
the
number

of
lines
left.
Sed
has
more
commands
that
make
grep
unnecessary.
And
grep
-c
can
replace
wc
-l.
I'll
discuss
how
you
can
duplicate
some
of
grep's
functionality
later.

Of
course
you
could
write
the
last
example
using
the
"-e"
option:

```
sed -e 's/^#.*//' f1 f2 f3 | grep -v '^$' | wc -l
```

There
are
two
other
options
to
sed.

sed
-n:
no
printing

The
"-n"
option
will
not
print
anything
unless

```
an
explicit
request
to
print
is
found.
I
mentioned
the
"/p"
flag
to
the
substitute
command
as
one
way
to
turn
printing
back
on.
Let
me
clarify
this.
The
command

sed 's/PATTERN/&/p' file

acts
like
the
cat
program
if
PATTERN
is
not
in
the
file:
e.g.
nothing
is
changed.
If
PATTERN
is
in
the
file,
then
each
line
that
has
this
is
printed
twice.
Add
the
"-n"
option
and
```

the
example
acts
like
grep:

`sed -n 's/PATTERN/&/p' file`

Nothing
is
printed,
except
those
lines
with
PATTERN
included.

The
long
argument
of
the
`-n`
command
is
either

`sed --quiet 's/PATTERN/&/p' file`

or

`sed --silent 's/PATTERN/&/p' file`

Using 'sed /pattern/'

Sed
has
the
ability
to
specify
which
lines
are
to
be
examined
and/or
modified,
by
specifying
addresses
before
the
command.
I
will
just
describe
the
simplest
version
for
now
-

the
/PATTERN/
address.
When
used,
only
lines
that
match
the
pattern
are
given
the
command
after
the
address.
Briefly,
when
used
with
the
/p
flag,
matching
lines
are
printed
twice:

sed '/PATTERN/p' file

And
of
course
PATTERN
is
any
regular
expression.

Please
note
that
if
you
do
not
include
a
command,
such
as
the
"p"
for
print,
you
will
get
an
error.
When
I
type

```
echo abc | sed '/a/'
```

I

```
get
the
error

sed: -e expression #1, char 3: missing command
```

Also,
you
don't
need
to,
but
I
recommend
that
you
place
a
space
after
the
pattern
and
the
command.
This
will
help
you
distinguish
between
flags
that
modify
the
pattern
matching,
and
commands
to
execute
after
the
pattern
is
matched.
Therefore
I
recommend
this
style:

```
sed '/PATTERN/ p' file
```

Using
'sed
-n
/pattern/p'
to
duplicate
the
function
of

grep

If
you
want
to
duplicate
the
functionality
of
grep,
combine
the
-n
(noprnt)
option
with
the
/p
print
flag:

```
sed -n '/PATTERN/p' file
```

sed

-f scriptname

If
you
have
a
large
number
of
sed
commands,
you
can
put
them
into
a
file
and
use

```
sed -f sedscrip <old >new
```

where
sedscrip
could
look
like
this:

```
# sed comment - This script changes lower case vowels to upper case  
s/a/A/g  
s/e/E/g  
s/i/I/g  
s/o/O/g  
s/u/U/g
```

When
there
are
several
commands

in
one
file,
each
command
must
be
on
a
separate
line.

The
long
argument
version
is

```
sed --file=sedscript <old >new
```

Also
see

[here](#)
[on](#)
[writing](#)
[a](#)
[script](#)
[that](#)
[executes](#)
[sed](#)
[directly](#)

sed in shell scripts

If
you
have
many
commands
and
they
won't
fit
neatly
on
one
line,
you
can
break
up
the
line
using
a
backslash:

```
sed -e 's/a/A/g' \  
    -e 's/e/E/g' \  
    -e 's/i/I/g' \  
    -e 's/o/O/g' \  
    -e 's/u/U/g' <old >new
```

Quoting multiple sed lines in the C shell

You
can
have
a
large,
multi-line

sed
script

in
the
C

shell,
but
you
must

tell
the
C

shell
that
the
quote
is
continued
across
several
lines.

This
is
done
by
placing
a
backslash

at
the
end
of
each
line:

```
#!/bin/csh -f
sed 's/a/A/g \
s/e/E/g \
s/i/I/g \
s/o/O/g \
s/u/U/g' <old >new
```

Quoting multiple sed lines

in the Bourne shell

The
Bourne
shell
makes
this
easier
as
a
quote
can
cover
several
lines:

```
#!/bin/sh
sed '
s/a/A/g
s/e/E/g
s/i/I/g
s/o/O/g
s/u/U/g' <old >new
```

sed -V

The
-V
option
will
print
the
version
of
sed
you
are
using.
The
long
argument
of
the
command
is

```
sed --version
```

sed -h

The
-h
option
will
print
a
summary
of
the

sed
commands.
The
long
argument
of
the
command
is

sed --help

A sed interpreter script

Another
way
of
executing
sed
is
to
use
an
interpreter
script.
Create
a
file
that
contains:

```
#!/bin/sed
-f
s/a/A/g
s/e/E/g
s/i/I/g
s/o/O/g
s/u/U/g
```

Click
here
to
get
file:

CapVowel.sed

If
this
script
was
stored
in
a
file
with
the
name
"CapVowel"
and
was
executable,
you
could

use
it
with
the
simple
command:

```
CapVowel <old >new
```

Comments

Sed
comments
are
lines
where
the
first
non-white
character
is
a
"#."
On
many
systems,
sed
can
have
only
one
comment,
and
it
must
be
the
first
line
of
the
script.
On
the
Sun
(1988
when
I
wrote
this),
you
can
have
several
comment
lines
anywhere
in
the
script.
Modern
versions
of
Sed
support
this.
If
the

first
line
contains
exactly
"#n"
then
this
does
the
same
thing
as
the
"-n"
option:
turning
off
printing
by
default.
This
could
not
done
with
a
sed
interpreter
script,
because
the
first
line
must
start
with
"#!/bin/sed
-f"
as
I
think
"#!/bin/sed
-nf"
generated
an
error.
It
worked
when
I
first
wrote
this
(2008).
Note
that
"#!/bin/sed
-fn"
does
not
work
because
sed
thinks
the
filename
of
the
script

```
is
"n".
However,
#!/bin/sed -nf"
does
work.
```

Passing **arguments** **into** **a** **sed** **script**

Passing
a
word
into
a
shell
script
that
calls
sed
is
easy
if
you
remembered
my
tutorial
on
the
UNIX
quoting
mechanism.

To
review,
you
use
the
single
quotes
to
turn
quoting
on
and
off.

A
simple
shell
script
that
uses
sed
to
emulate
grep
is:

```
#!/bin/sh
sed
```

```
-n  
's/'$1'/&/p'
```

However

there
is
a
problem
with
this
script.
If
you
have
a
space
as
an
argument,
the
script
would
cause
a
syntax
error
A
better
version
would
protect
from
this
happening:

```
#!/bin/sh  
sed -n 's/'"$1"'/&/p'
```

Click
here
to
get
file:

[**sedgrep.sed**](#)

If
this
was
stored
in
a
file
called
sedgrep,
you
could
type

```
sedgrep '[A-Z][A-Z]' <file
```

This
would
allow
sed
to
act
as
the

grep
command.

Using sed in a shell here-is document

You
can
use
sed
to
prompt
the
user
for
some
parameters
and
then
create
a
file
with
those
parameters
filled
in.
You
could
create
a
file
with
dummy
values
placed
inside
it,
and
use
sed
to
change
those
dummy
values.
A
simpler
way
is
to
use
the
"here
is"
document,
which
uses
part

```
of
the
shell
script
as
if
it
were
standard
input:

#!/bin/sh
echo -n 'what is the value? '
read value
sed 's/XYZ/'$value'/' <<EOF
The value is XYZ
EOF

When
executed,
the
script
says:

what is the value?

If
you
type
in
"123,"
the
next
line
will
be:

The value is 123

I
admit
this
is
a
contrived
example.
"Here
is"
documents
can
have
values
evaluated
without
the
use
of
sed.
This
example
does
the
same
thing:

#!/bin/sh
echo -n 'what is the value? '
read value
cat <<EOF
The value is $value
EOF
```

However,
combining
"here
is"
documents
with
sed
can
be
useful
for
some
complex
cases.
Note
that

```
sed  
's/XYZ  
/'$value'/'  
<<EOF
```

will
give
a
syntax
error
if
the
user
types
an
answer
that
contains
a
space,
like
"a
b
c".
Better
form
would
be
to
put
double
quotes
around
the
evaluation
of
the
value:

```
#!/bin/sh  
echo -n 'what is the value? '  
read value  
sed 's/XYZ/'"$value"'/' <<EOF  
The value is XYZ  
EOF
```

I
covered
this
in
my
tutorial
on

quotation marks .

Click
here
to
get
file:

[sed_hereis.sed](#)

Multiple commands and order of execution

As
we
explore
more
of
the
commands
of
sed,
the
commands
will
become
complex,
and
the
actual
sequence
can
be
confusing.
It's
really
quite
simple.
Each
line
is
read
in.
Each
command,
in
order
specified
by
the
user,
has
a
chance
to
operate
on
the
input
line.
After

the
substitutions
are
made,
the
next
command
has
a
chance
to
operate
on
the
same
line,
which
may
have
been
modified
by
earlier
commands.
If
you
ever
have
a
question,
the
best
way
to
learn
what
will
happen
is
to
create
a
small
example.
If
a
complex
command
doesn't
work,
make
it
simpler.
If
you
are
having
problems
getting
a
complex
script
working,
break
it
up
into
two
smaller

scripts
and
pipe
the
two
scripts
together.

Addresses **and** **Ranges** **of** **Text**

You
have
only
learned
one
command,
and
you
can
see
how
powerful
sed
is.
However,
all
it
is
doing
is
a
grep
and
substitute.
That
is,
the
substitute
command
is
treating
each
line
by
itself,
without
caring
about
nearby
lines.
What
would
be
useful
is
the
ability
to
restrict
the
operation

to
certain
lines.
Some
useful
restrictions
might
be:

Specifying
a
line
by
its
number.
Specifying
a
range
of
lines
by
number.
All
lines
containing
a
pattern.
All
lines
from
the
beginning
of
a
file
to
a
regular
expression
All
lines
from
a
regular
expression
to
the
end
of
the
file.
All
lines
between
two
regular
expressions.

Sed
can
do
all
that
and
more.
Every

command
in
sed
can
be
preceded
by
an
address,
range
or
restriction
like
the
above
examples.
The
restriction
or
address
immediately
precedes
the
command:

restriction
command

Restricting **to** **a** **line** **number**

The
simplest
restriction
is
a
line
number.
If
you
wanted
to
delete
the
first
number
on
line
3,
just
add
a
"3"
before
the
command:

```
sed '3 s/[0-9][0-9]*//' <file >new
```

Patterns

Many

UNIX
utilities
like
vi
and
more
use
a
slash
to
search
for
a
regular
expression.
Sed
uses
the
same
convention,
provided
you
terminate
the
expression
with
a
slash.
To
delete
the
first
number
on
all
lines
that
start
with
a
"#,"
use:

`sed '/^#/ s/[0-9][0-9]*//'`

I
placed
a
space
after
the
"/expression/"
so
it
is
easier
to
read.
It
isn't
necessary,
but
without
it
the
command
is
harder
to

fathom.
Sed
does
provide
a
few
extra
options
when
specifying
regular
expressions.
But
I'll
discuss
those
later.
If
the
expression
starts
with
a
backslash,
the
next
character
is
the
delimiter.
To
use
a
comma
instead
of
a
slash,
use:

`sed '\,^#, s/[0-9][0-9]*//'`

The
main
advantage
of
this
feature
is
searching
for
slashes.
Suppose
you
wanted
to
search
for
the
string
"/usr/local/bin"
and
you
wanted
to
change
it
for
"/common

```
/all/bin."
You
could
use
the
backslash
to
escape
the
slash:

sed '/usr/local/bin/ s/\/usr/local/\/common/all/'

It
would
be
easier
to
follow
if
you
used
an
underline
instead
of
a
slash
as
a
search.
This
example
uses
the
underline
in
both
the
search
command
and
the
substitute
command:

sed '\_usr/local/bin_ s/_usr/local_/common/all_'

This
illustrates
why
sed
scripts
get
the
reputation
for
obscurity.
I
could
be
perverse
and
show
you
the
example
that
will
search
```

```
for
all
lines
that
start
with
a
"g,"
and
change
each
"g"
on
that
line
to
an
"s:"

sed '/^g/s/g/s/g'
```

```
Adding
a
space
and
using
an
underscore
after
the
substitute
command
makes
this
much
easier
to
read:

sed '/^g/ s_g_s_g'
```

```
Er,
I
take
that
back.
It's
hopeless.
There
is
a
lesson
here:
Use
comments
liberally
in
a
sed
script.
You
may
have
to
remove
the
comments
to
run
the
```

script
under
a
different
(older)
operating
system,
but
you
now
know
how
to
write
a
sed
script
to
do
that
very
easily!
Comments
are
a
Good
Thing.
You
may
have
understood
the
script
perfectly
when
you
wrote
it.
But
six
months
from
now
it
could
look
like
modem
noise.
And
if
you
don't
understand
that
reference,
imagine
an
8-month-old
child
typing
on
a
computer.

Ranges
by

line number

You
can
specify
a
range
on
line
numbers
by
inserting
a
comma
between
the
numbers.

To
restrict
a
substitution
to
the
first
100
lines,
you
can
use:

```
sed '1,100 s/A/a/'
```

If
you
know
exactly
how
many
lines
are
in
a
file,
you
can
explicitly
state
that
number
to
perform
the
substitution
on
the
rest
of
the
file.

In
this
case,
assume
you
used
`wc`
to

```
find
out
there
are
532
lines
in
the
file:

sed '101,532 s/A/a/'

An
easier
way
is
to
use
the
special
character
"$,"
which
means
the
last
line
in
the
file.

sed '101,$ s/A/a/'

The
"$"
is
one
of
those
conventions
that
mean
"last"
in
utilities
like
cat
-e,
vi,
and
ed.
"cat
-e"
Line
numbers
are
cumulative
if
several
files
are
edited.
That
is,

sed '200,300 s/A/a/' f1 f2 f3 >new

is
the
same
```

```
as
cat f1 f2 f3 | sed '200,300 s/A/a/' >new
```

Ranges by patterns

You can specify two regular expressions as the range. Assuming a "#" starts a comment, you can search for a keyword, remove all comments until you see the second keyword. In this case the two keywords are "start" and "stop:"

```
sed '/start/,/stop/ s/#.*//'
```

The first pattern turns on a flag that tells *sed* to perform the substitute command on

every
line.
The
second
pattern
turns
off
the
flag.
If
the
"start"
and
"stop"
pattern
occurs
twice,
the
substitution
is
done
both
times.
If
the
"stop"
pattern
is
missing,
the
flag
is
never
turned
off,
and
the
substitution
will
be
performed
on
every
line
until
the
end
of
the
file.

You
should
know
that
if
the
"start"
pattern
is
found,
the
substitution
occurs
on
the
same
line
that

```
contains  
"start."  
This  
turns  
on  
a  
switch,  
which  
is  
line  
oriented.  
That  
is,  
the  
next  
line  
is  
read  
and  
the  
substitute  
command  
is  
checked.  
If  
it  
contains  
"stop"  
the  
switch  
is  
turned  
off.  
Switches  
are  
line  
oriented,  
and  
not  
word  
oriented.  
  
You  
can  
combine  
line  
numbers  
and  
regular  
expressions.  
This  
example  
will  
remove  
comments  
from  
the  
beginning  
of  
the  
file  
until  
it  
finds  
the  
keyword  
"start:"  
  
sed -e '1,/start/ s/#.*//'
```

This
example
will
remove
comments
everywhere
except
the
lines
between
the
two
keywords:

```
sed -e '1,/start/ s/#.*//' -e '/stop/, $ s/#.*//'
```

The
last
example
has
a
range
that
overlaps
the
"/start
/,/stop/"
range,
as
both
ranges
operate
on
the
lines
that
contain
the
keywords.

I
will
show
you
later
how
to
restrict
a
command
up
to,
but
not
including
the
line
containing
the
specified
pattern.
It
is
in

Operating
in
a
pattern
range
except

for the patterns

But
I
have
to
cover
some
more
basic
principles.

Before
I
start
discussing
the
various
commands,
I
should
explain
that
some
commands
cannot
operate
on
a
range
of
lines.

I
will
let
you
know
when
I
mention
the
commands.

In
this
next
section

I
will
describe
three
commands,
one
of
which
cannot
operate
on
a
range.

Delete with d

Using

ranges
can
be
confusing,
so
you
should
expect
to
do
some
experimentation
when
you
are
trying
out
a
new
script.
A
useful
command
deletes
every
line
that
matches
the
restriction:
"d."
If
you
want
to
look
at
the
first
10
lines
of
a
file,
you
can
use:

```
sed '11,$ d' <file
```

which
is
similar
in
function
to
the
head
command.
If
you
want
to
chop
off
the
header
of
a


```
mail
message,
which
is
everything
up
to
the
first
blank
line,
use:

sed '1,/^\$/ d' <file
```

```
You
can
duplicate
the
function
of
the
tail
command,
assuming
you
know
the
length
of
a
file.
Wc
can
count
the
lines,
and
expr
can
subtract
10
from
the
number
of
lines.
A
Bourne
shell
script
to
look
at
the
last
10
lines
of
a
file
might
look
like
this:
```

```
#!/bin/sh
#print
last
```

```
10
lines
of
file
#
First
argument
is
the
filename
lines=`wc
-l
$1
|
awk
'{print
$1}'
`
start=`expr
$lines
-
10`
sed
"1,$start
d"
$1
```

Click
here
to
get
file:
[sed_tail.sh](#)

The
range
for
deletions
can
be
regular
expressions
pairs
to
mark
the
begin
and
end
of
the
operation.
Or
it
can
be
a
single
regular
expression.
Deleting
all
lines
that
start
with
a
"#"
is

easy:

```
sed '/^#/ d'
```

Removing
comments
and
blank
lines
takes
two
commands.
The
first
removes
every
character
from
the
"#"
to
the
end
of
the
line,
and
the
second
deletes
all
blank
lines:

```
sed -e 's/#.*//' -e '/^$/ d'
```

A
third
one
should
be
added
to
remove
all
blanks
and
tabs
immediately
before
the
end
of
line:

```
sed -e 's/#.*//' -e 's/[ ^I]*$//' -e '/^$/ d'
```

The
character
"^I"
is
a
CTRL-I
or
tab
character.
You
would
have
to

explicitly
type
in
the
tab.
Note
the
order
of
operations
above,
which
is
in
that
order
for
a
very
good
reason.
Comments
might
start
in
the
middle
of
a
line,
with
white
space
characters
before
them.
Therefore
comments
are
first
removed
from
a
line,
potentially
leaving
white
space
characters
that
were
before
the
comment.
The
second
command
removes
all
trailing
blanks,
so
that
lines
that
are
now
blank
are

converted
to
empty
lines.
The
last
command
deletes
empty
lines.
Together,
the
three
commands
remove
all
lines
containing
only
comments,
tabs
or
spaces.

This
demonstrates
the
pattern
space
sed
uses
to
operate
on
a
line.
The
actual
operation
sed
uses
is:

Copy
the
input
line
into
the
pattern
space.
Apply
the
first
sed
command
on
the
pattern
space,
if
the
address
restriction
is
true.
Repeat

with
the
next
sed
expression,
again
operating
on
the
pattern
space.
When
the
last
operation
is
performed,
write
out
the
pattern
space
and
read
in
the
next
line
from
the
input
file.

Printing **with** **p**

Another
useful
command
is
the
print
command:
"p."
If
sed
wasn't
started
with
an
"-n"
option,
the
"p"
command
will
duplicate
the
input.
The
command

`sed 'p'`

will
duplicate
every
line.
If
you
wanted
to
double
every
empty
line,
use:

```
sed '/^$/ p'
```

Adding
the
"-n"
option
turns
off
printing
unless
you
request
it.
Another
way
of
duplicating
head's
functionality
is
to
print
only
the
lines
you
want.
This
example
prints
the
first
10
lines:

```
sed -n '1,10 p' <file
```

Sed
can
act
like
grep
by
combining
the
print
operator
to
function
on
all
lines
that
match
a
regular

expression:

```
sed -n '/match/ p'
```

which
is
the
same
as:

```
grep match
```

Reversing the restriction with !

Sometimes
you
need
to
perform
an
action
on
every
line
except
those
that
match
a
regular
expression,
or
those
outside
of
a
range
of
addresses.
The
"!"
character,
which
often
means
not
in
UNIX
utilities,
inverts
the
address
restriction.
You
remember
that

```
sed -n '/match/ p'
```

acts
like
the

grep
command.
The
"-v"
option
to
grep
prints
all
lines
that
don't
contain
the
pattern.
Sed
can
do
this
with

`sed -n '/match/ !p' </tmp/b`

Relationships **between** **d,** **p,** **and** **!**

As
you
may
have
noticed,
there
are
often
several
ways
to
solve
the
same
problem
with
sed.
This
is
because
print
and
delete
are
opposite
functions,
and
it
appears
that
"!p"
is
similar
to
"d,"

while
"!d"
is
similar
to
"p."
I
wanted
to
test
this,
so
I
created
a
20
line
file,
and
tried
every
different
combination.
The
following
table,
which
shows
the
results,
demonstrates
the
difference:

Relations between d, p, and !			
Sed	Range	Command	Results
sed -n	1,10	p	Print first 10 lines
sed -n	11,\$!p	Print first 10 lines
sed	1,10	!d	Print first 10 lines
sed	11,\$	d	Print first 10 lines
sed -n	1,10	!p	Print last 10 lines
sed -n	11,\$	p	Print last 10 lines
sed	1,10	d	Print last 10 lines
sed	11,\$!d	Print last 10 lines
sed -n	1,10	d	Nothing printed
sed -n	1,10	!d	Nothing printed
sed -n	11,\$	d	Nothing printed
sed -n	11,\$!d	Nothing printed

sed	1,10	p	Print first 10 lines twice, then next 10 lines once
sed	11,\$!p	Print first 10 lines twice, then last 10 lines once
sed	1,10	!p	Print first 10 lines once, then last 10 lines twice
sed	11,\$	p	Print first 10 lines once, then last 10 lines twice

This table shows that the following commands are identical:

```
sed -n '1,10 p'  
sed -n '11,$ !p'  
sed '1,10 !d'  
sed '11,$ d'
```

It also shows that the "!" command "inverts" the address range, operating on the other lines.

The q or quit command

There
is
one
more
simple
command
that
can
restrict
the
changes
to
a
set
of
lines.
It
is
the
"q"
command:
quit.
The
third
way
to
duplicate
the
head
command
is:

```
sed '11 q'
```

which
quits
when
the
eleventh
line
is
reached.
This
command
is
most
useful
when
you
wish
to
abort
the
editing
after
some
condition
is
reached.

The
"q"
command
is
the
one
command
that
does

not
take
a
range
of
addresses.
Obviously
the
command

```
sed '1,10 q'
```

cannot
quit
10
times.
Instead

```
sed '1 q'
```

or

```
sed '10 q'
```

is
correct.

Grouping with { and }

The
curly
braces,
"{"
and
"}",
are
used
to
group
the
commands.

Hardly
worth
the
buildup.
All
that
prose
and
the
solution
is
just
matching
squiggles.
Well,
there
is
one
complication.
Since
each

`sed`
command
must
start
on
its
own
line,
the
curly
braces
and
the
nested
`sed`
commands
must
be
on
separate
lines.

Previously,
I
showed
you
how
to
remove
comments
starting
with
a
"`#.`".
If
you
wanted
to
restrict
the
removal
to
lines
between
special
"`begin`"
and
"`end`"
key
words,
you
could
use:

```
#!/bin/sh
# This is a Bourne shell script that removes #-type comments
# between 'begin' and 'end' words.
sed -n '
    /begin/,/end/ {
        s/#.*//
        s/[ ^I]*$//
        /^$/ d
        p
    }
'
```

Click
here
to

get
file:
[sed_begin_end.sh](#)

These
braces
can
be
nested,
which
allow
you
to
combine
address
ranges.
You
could
perform
the
same
action
as
before,
but
limit
the
change
to
the
first
100
lines:

```
#!/bin/sh
# This is a Bourne shell script that removes #-type comments
# between 'begin' and 'end' words.
sed -n '
    1,100 {
        /begin/,/end/ {
            s/#.*//
            s/[ ^I]*$//
            /^$/ d
            p
        }
    }
'
```

Click
here
to
get
file:
[sed_begin_end1.sh](#)

You
can
place
a
"!"
before
a
set
of
curly
braces.
This
inverts
the
address,
which

removes
comments
from
all
lines
except
those
between
the
two
reserved
words:

```
#!/bin/sh
sed '
    /begin/,/end/ !{
        s/#.*//
        s/[ ^I]*$//
        /^$/ d
    }
'
```

Click
here
to
get
file:

[sed_begin_end2.sh](#)

Operating in a pattern range except for the patterns

You
may
remember
that
I
mentioned
you
can
do
a
substitute
on
a
pattern
range,
like
changing
"old"
to
"new"
between
a


```
begin/end
pattern:

#!/bin/sh
sed '
    /begin/,/end/ s/old/new/
'
```

Another way to write this is to use the curly braces for grouping:

```
#!/bin/sh
sed '
    /begin/,/end/ {
        s/old/new/
    }
'
```

I think this makes the code clearer to understand, and easier to modify, as you will see below.

If you did not want to make any changes where the word "begin" occurred, you could simply add a new condition to skip

```
over
that
line:

#!/bin/sh
sed '
    /begin/,/end/ {
        /begin/n # skip over the line that has "begin" on it
        s/old/new/
    }
'
```

However, skipping over the line that has "end" is trickier. If you use the same method you used for "begin" then the sed engine will not see the "end" to stop the range - it skips over that as well. The solution is to do a substitute on all lines that don't have the "end" by using

```
#!/bin/sh
sed '
    /begin/,/end/ {
        /begin/n # skip over the line that has "begin" on it
        /end/ !{
            s/old/new/
        }
    }
'
```

Writing a file with the 'w' command

You may remember that the substitute command can write to a file. Here again is the example that will only write lines that start with an even number (and followed by a space):

```
sed -n 's/^[0-9]*[02468] /&/w even' <file
```

I used the "&" in the replacement part of the substitution command

so
that
the
line
would
not
be
changed.
A
simpler
example
is
to
use
the
"w"
command,
which
has
the
same
syntax
as
the
"w"
flag
in
the
substitute
command:

`sed -n '/^[0-9]*[02468]/ w even' <file`

Remember

-
only
one
space
must
follow
the
command.
Anything
else
will
be
considered
part
of
the
file
name.
The
"w"
command
also
has
the
same
limitation
as
the
"w"
flag:
only
10
files
can
be

opened
in
sed.

Reading in a file with the 'r' command

There
is
also
a
command
for
reading
files.
The
command

```
sed '$r end' <in>out
```

will
append
the
file
"end"
at
the
end
of
the
file
(address
"\$")."
The
following
will
insert
a
file
after
the
line
with
the
word
"INCLUDE:"

```
sed '/INCLUDE/ r file' <in >out
```

You
can
use
the
curly
braces
to
delete
the

line
having
the
"INCLUDE"
command
on
it:

```
#!/bin/sh  
sed '/INCLUDE/ {  
    r file  
    d  
}'
```

Click
here
to
get
file:

[sed_include.sh](#)

The
order
of
the
delete
command
"d"
and
the
read
file
command
"r"
is
important.
Change
the
order
and
it
will
not
work.
There
are
two
subtle
actions
that
prevent
this
from
working.
The
first
is
the
"r"
command
writes
the
file
to
the
output
stream.
The

file
is
not
inserted
into
the
pattern
space,
and
therefore
cannot
be
modified
by
any
command.
Therefore
the
delete
command
does
not
affect
the
data
read
from
the
file.

The
other
subtlety
is
the
"d"
command
deletes
the
current
data
in
the
pattern
space.
Once
all
of
the
data
is
deleted,
it
does
make
sense
that
no
other
action
will
be
attempted.
Therefore
a
"d"
command
executed
in

a
curly
brace
also
aborts
all
further
actions.
As
an
example,
the
substitute
command
below
is
never
executed:

```
#!/bin/sh
# this example is WRONG
sed -e '1 {
    d
    s/.*/ /
}'
```

Click
here
to
get
file:
[sed_bad_example.sh](#)

The
earlier
example
is
a
crude
version
of
the
C
preprocessor
program.
The
file
that
is
included
has
a
predetermined
name.
It
would
be
nice
if
sed
allowed
a
variable
(e.g
"1"
)
instead
of

a
fixed
file
name.
Alas,
sed
doesn't
have
this
ability.
You
could
work
around
this
limitation
by
creating
sed
commands
on
the
fly,
or
by
using
shell
quotes
to
pass
variables
into
the
sed
script.
Suppose
you
wanted
to
create
a
command
that
would
include
a
file
like
cpp,
but
the
filename
is
an
argument
to
the
script.
An
example
of
this
script
is:

```
% include 'sys/param.h' <file.c >file.c.new
```

A
shell

script
to
do
this
would
be:

```
#!/bin/sh
# watch out for a '/' in the parameter
# use alternate search delimiter
sed -e '\_#INCLUDE <"$1">_{
    r "$1"
    d
}'
```

Let
me
elaborate.
If
you
had
a
file
that
contains

```
Test first file
#INCLUDE <file1>
Test second file
#INCLUDE <file2>
```

you
could
use
the
command

```
sed_include1.sh file1<input|sed_include1.sh file2
```

to
include
the
specified
files.

Click
here
to
get
file:

[sed_include1.sh](#)

SunOS and the # Comment Command

As
we
dig
deeper
into
sed,
comments

will
make
the
commands
easier
to
follow.
The
older
versions
of
sed
only
allow
one
line
as
a
comment,
and
it
must
be
the
first
line.
SunOS
(and
GNU's
sed)
allows
more
than
one
comment,
and
these
comments
don't
have
to
be
first.
The
last
example
could
be:

```
#!/bin/sh
# watch out for a '/' in the parameter
# use alternate search delimiter
sed -e '\_#INCLUDE <"$1">_{
    # read the file
    r "$1"

    # delete any characters in the pattern space
    # and read the next line in
    d
}'
```

Click
here
to
get
file:

[**sed_include2.sh**](#)

Adding, **Changing,** **Inserting** **new** **lines**

Sed
has
three
commands
used
to
add
new
lines
to
the
output
stream.
Because
an
entire
line
is
added,
the
new
line
is
on
a
line
by
itself
to
emphasize
this.
There
is
no
option,
an
entire
line
is
used,
and
it
must
be
on
its
own
line.
If
you
are
familiar
with
many
UNIX
utilities,
you
would
expect

sed
to
use
a
similar
convention:
lines
are
continued
by
ending
the
previous
line
with
a
"\".
The
syntax
to
these
commands
is
finicky,
like
the
"r"
and
"w"
commands.

Append

a line with 'a'

The
"a"
command
appends
a
line
after
the
range
or
pattern.
This
example
will
add
a
line
after
every
line
with
"WORD:"

```
#!/bin/sh
sed '
/WORD/ a\
Add this line after every line with WORD
'
```

Click
here
to
get
file:
[sed add line after word.sh](#)

You
could
eliminate
two
lines
in
the
shell
script
if
you
wish:

```
#!/bin/sh
sed '/WORD/ a\
Add this line after every line with WORD'
```

Click
here
to
get
file:
[sed add line after word1.sh](#)

I
prefer
the
first
form
because
it's
easier
to
add
a
new
command
by
adding
a
new
line
and
because
the
intent
is
clearer.
There
must
not
be
a
space
after
the
"\".

Insert

a line with 'i'

You
can
insert
a
new
line
before
the
pattern
with
the
"i"
command:

```
#!/bin/sh
sed '
/WORD/ i\
Add this line before every line with WORD
'
```

Click
here
to
get
file:

[sed_add_line_before_word.sh](#)

Change a line with 'c'

You
can
change
the
current
line
with
a
new
line.

```
#!/bin/sh
sed '
/WORD/ c\
Replace the current line with the line
'
```

Click
here
to
get
file:

[sed_change_line.sh](#)

A
"d"
command
followed
by
a
"a"
command
won't
work,
as
I
discussed
earlier.
The
"d"
command
would
terminate
the
current
actions.
You
can
combine
all
three
actions
using
curly
braces:

```
#!/bin/sh
sed '
/WORD/ {
i\
Add this line before
a\
Add this line after
c\
Change the line to this one
}'
```

Click
here
to
get
file:

[sed_insert_append_change.sh](#)

[Leading tabs and spaces in a sed script](#)

Sed
ignores

leading
tabs
and
spaces
in
all
commands.
However
these
white
space
characters
may
or
may
not
be
ignored
if
they
start
the
text
following
a
"a,"
"c"
or
";"
command.
In
SunOS,
both
"features"
are
available.
The
Berkeley
(and
Linux)
style
sed
is
in
/usr/bin,
and
the
AT&T
version
(System
V)
is
in
/usr/5bin/.
To
elaborate,
the
/usr/bin/sed
command
retains
white
space,
while
the
/usr/5bin/sed
strips
off
leading

```
spaces.
If
you
want
to
keep
leading
spaces,
and
not
care
about
which
version
of
sed
you
are
using,
put
a
"\
as
the
first
character
of
the
line:

#!/bin/sh
sed '
    a\
    This line starts with a tab
'
```

Adding more than one line

```
All
three
commands
will
allow
you
to
add
more
than
one
line.
Just
end
each
line
with
a
"\:"

#!/bin/sh
sed '
/WORD/ a\
Add this line\
This line\
```

And this line

Adding lines and the pattern space

I have mentioned the pattern space before. Most commands operate on the pattern space, and subsequent commands may act on the results of the last modification. The three previous commands, like the read file command, add the new lines to the output stream, bypassing the pattern space.

Address ranges and the

above commands

You
may
remember
that
earlier
I
warned
you
that
some
commands
can
take
a
range
of
lines,
and
others
cannot.
To
be
precise,
the
commands
"a,"
"i,"
"r,"
and
"q"
will
not
take
a
range
like
"1,100"
or
"/begin
/,/end/."
The
documentation
states
that
the
read
command
can
take
a
range,
but
I
got
an
error
when
I
tried
this.
The
"c"
or
change

command
allows
this,
and
it
will
let
you
change
several
lines
into
one:

```
#!/bin/sh
sed '
/begin/,/end/ c\
***DELETED***
'
```

If
you
need
to
do
this,
you
can
use
the
curly
braces,
as
that
will
let
you
perform
the
operation
on
every
line:

```
#!/bin/sh
# add a blank line after every line
sed '1,$ {
    a\
}'
```

Multi-Line Patterns

Most
UNIX
utilities
are
line
oriented.
Regular
expressions
are
line
oriented.
Searching
for
patterns

that
covers
more
than
one
line
is
not
an
easy
task.
(Hint:
It
will
be
very
shortly.)

Sed
reads
in
a
line
of
text,
performs
commands
which
may
modify
the
line,
and
outputs
modification
if
desired.
The
main
loop
of
a
sed
script
looks
like
this:

The
next
line
is
read
from
the
input
file
and
places
it
in
the
pattern
space.
If
the
end
of

file
is
found,
and
if
there
are
additional
files
to
read,
the
current
file
is
closed,
the
next
file
is
opened,
and
the
first
line
of
the
new
file
is
placed
into
the
pattern
space.
The
line
count
is
incremented
by
one.
Opening
a
new
file
does
not
reset
this
number.
Each
sed
command
is
examined.
If
there
is
a
restriction
placed
on
the
command,
and

the
current
line
in
the
pattern
space
meets
that
restriction,
the
command
is
executed.
Some
commands,
like
"n"
or
"d"
cause
sed
to
go
to
the
top
of
the
loop.
The
"q"
command
causes
sed
to
stop.
Otherwise
the
next
command
is
examined.
After
all
of
the
commands
are
examined,
the
pattern
space
is
output
unless
sed
has
the
optional
"-n"
argument.

The
restriction
before

the
command
determines
if
the
command
is
executed.
If
the
restriction
is
a
pattern,
and
the
operation
is
the
delete
command,
then
the
following
will
delete
all
lines
that
have
the
pattern:

```
/PATTERN/ d
```

If
the
restriction
is
a
pair
of
numbers,
then
the
deletion
will
happen
if
the
line
number
is
equal
to
the
first
number
or
greater
than
the
first
number
and
less
than
or
equal

```
to
the
last
number:
10,20 d

If
the
restriction
is
a
pair
of
patterns,
there
is
a
variable
that
is
kept
for
each
of
these
pairs.
If
the
variable
is
false
and
the
first
pattern
is
found,
the
variable
is
made
true.
If
the
variable
is
true,
the
command
is
executed.
If
the
variable
is
true,
and
the
last
pattern
is
on
the
line,
after
the
command
is
```

executed
the
variable
is
turned
off:

```
/begin/,/end/ d
```

Whew!
That
was
a
mouthful.
If
you
have
read
carefully
up
to
here,
you
should
have
breezed
through
this.
You
may
want
to
refer
back,
because
I
covered
several
subtle
points.
My
choice
of
words
was
deliberate.
It
covers
some
unusual
cases,
like:

```
# what happens if the second number  
# is less than the first number?  
sed -n '20,1 p' file
```

and

```
# generate a 10 line file with line numbers  
# and see what happens when two patterns overlap  
yes | head -10 | cat -n | \  
sed -n -e '/1/,/7/ p' -e '/5/,/9/ p'
```

Enough
mental
punishment.
Here
is
another
review,

this
time
in
a
table
format.
Assume
the
input
file
contains
the
following
lines:

AB
CD
EF
GH
IJ

When
sed
starts
up,
the
first
line
is
placed
in
the
pattern
space.
The
next
line
is
"CD."
The
operations
of
the
"n,"
"d,"
and
"p"
commands
can
be
summarized
as:

Pattern Space	Next Input	Command	Output	New Pattern Space	New Text Input
AB	CD	n	<default>	CD	EF
AB	CD	d	-	CD	EF
AB	CD	p	AB	CD	EF

The
"n"
command
may
or
may
not
generate
output
depending

upon
the
existence
of
the
"-n"
flag.

That
review
is
a
little
easier
to
follow,
isn't
it?
Before
I
jump
into
multi-line
patterns,
I
wanted
to
cover
three
more
commands:

Print **line** **number** **with** **=**

The
"=" command
prints
the
current
line
number
to
standard
output.
One
way
to
find
out
the
line
numbers
that
contain
a
pattern
is
to
use:

```
# add line numbers first,
# then use grep,
# then just print the number
cat -n file | grep 'PATTERN' | awk '{print $1}'
```

The
sed
solution
is:

```
sed -n '/PATTERN/ =' file
```

Earlier
I
used
the
following
to
find
the
number
of
lines
in
a
file

```
#!/bin/sh
lines=`wc -l file | awk '{print $1}'`
```

Using
the
"=" command
can
simplify
this:

```
#!/bin/sh
lines=`sed -n '$=' file`
```

The
"=" command
only
accepts
one
address,
so
if
you
want
to
print
the
number
for
a
range
of
lines,
you
must
use
the
curly
braces:

```
#!/bin/sh
# Just print the line numbers
sed -n '/begin/,/end/ {
=
}
```

```
d  
}' file
```

Since the "=" command only prints to standard output, you cannot print the line number on the same line as the pattern. You need to edit multi-line patterns to do this.

Transform with y

If you wanted to change a word from lower case to upper case, you could write
26
character substitutions, converting "a" to "A," etc.
Sed has a

command
that
operates
like
the
tr
program.
It
is
called
the
"y"
command.
For
instance,
to
change
the
letters
"a"
through
"f"
into
their
upper
case
form,
use:

```
sed 'y/abcdef/ABCDEF/' file
```

Here's
a
sed
example
that
converts
all
uppercase
letters
to
lowercase
letters,
like
the
tr
command:

```
sed 'y/ABCDEFGHIJKLMNPOQRSTUVWXYZ/abcdefghijklmnopqrstuvwxyz/' <uppercase >lowercase
```

If
you
wanted
to
convert
a
line
that
contained
a
hexadecimal
number
(e.g.
0x1aff)
to
upper
case
(0x1AFF),
you
could


```
use:
sed '/0x[0-9a-zA-Z]*/ y/abcdef/ABCDEF' file

This
works
fine
if
there
are
only
numbers
in
the
file.
If
you
wanted
to
change
the
second
word
in
a
line
to
upper
case,
and
you
are
using
classic
sed,
you
are
out
of
luck
-
unless
you
use
multi-line
editing.
(Hey
-
I
think
there
is
some
sort
of
theme
here!)

However,
GNU
sed
has
a
uppercase
and
lowercase
extension.
```

Displaying control characters with a !

The
"!"
command
prints
the
current
pattern
space.
It
is
therefore
useful
in
debugging
sed
scripts.
It
also
converts
unprintable
characters
into
printing
characters
by
outputting
the
value
in
octal
preceded
by
a
"\
character.
I
found
it
useful
to
print
out
the
current
pattern
space,
while
probing
the
subtleties
of
sed.

Working with

Multiple Lines

There
are
three
new
commands
used
in
multiple-line
patterns:

"N,"
"D,"
and
"p."

I
will
explain
their
relation
to
the
matching
"n,"
"d,"
and
"p"
single-line
commands.

The
"n"
command
will
print
out
the
current
pattern
space
(unless
the
"-n"
flag
is
used),
empty
the
current
pattern
space,
and
read
in
the
next
line
of
input.
The
"N"
command
does
not
print
out

the
current
pattern
space
and
does
not
empty
the
pattern
space.
It
reads
in
the
next
line,
but
appends
a
new
line
character
along
with
the
input
line
itself
to
the
pattern
space.

The
"d"
command
deletes
the
current
pattern
space,
reads
in
the
next
line,
puts
the
new
line
into
the
pattern
space,
and
aborts
the
current
command,
and
starts
execution
at
the
first
sed
command.
This

is
called
starting
a
new
"cycle."
The
"D"
command
deletes
the
first
portion
of
the
pattern
space,
up
to
the
new
line
character,
leaving
the
rest
of
the
pattern
alone.
Like
"d,"
it
stops
the
current
command
and
starts
the
command
cycle
over
again.
However,
it
will
not
print
the
current
pattern
space.
You
must
print
it
yourself,
a
step
earlier.
If
the
"D"
command
is
executed
with
a

group
of
other
commands
in
a
curly
brace,
commands
after
the
"D"
command
are
ignored.
The
next
group
of
sed
commands
is
executed,
unless
the
pattern
space
is
emptied.
If
this
happens,
the
cycle
is
started
from
the
top
and
a
new
line
is
read.

The
"p"
command
prints
the
entire
pattern
space.
The
"p"
command
only
prints
the
first
part
of
the
pattern
space,
up
to
the

```
NEWLINE
character.
Neither
the
"p"
nor
the
"p"
command
changes
the
patterns
space.

Some
examples
might
demonstrate
"N"
by
itself
isn't
very
useful.
the
filter

sed -e 'N'

doesn't
modify
the
input
stream.
Instead,
it
combines
the
first
and
second
line,
then
prints
them,
combines
the
third
and
fourth
line,
and
prints
them,
etc.
It
does
allow
you
to
use
a
new
"anchor"
character:
"\n."
This
matches
the
```

new
line
character
that
separates
multiple
lines
in
the
pattern
space.
If
you
wanted
to
search
for
a
line
that
ended
with
the
character
"#,"
and
append
the
next
line
to
it,
you
could
use

```
#!/bin/sh
sed '
# look for a "#" at the end of the line
/##$/ {
# Found one - now read in the next line
    N
# delete the "#" and the new line character,
    s/##\n//
}' file
```

You
could
search
for
two
lines
containing
"ONE"
and
"TWO"
and
only
print
out
the
two
consecutive
lines:

```
#!/bin/sh
sed -n '
/ONE/ {
# found "ONE" - read in next line
    N
# look for "TWO" on the second line
# and print if there.
```



```

        /\n.*TWO/ p
    }' file

The
next
example
would
delete
everything
between
"ONE"
and
"TWO:"

#!/bin/sh
sed '
/ONE/ {
# append a line
    N
# search for TWO on the second line
    /\n.*TWO/ {
# found it - now edit making one line
        s/ONE.*\n.*TWO/ONE TWO/
    }
}' file

```

Matching three lines with sed

You
can
match
multiple
lines
in
searches.

Here
is
a
way
to
look
for
the
string
"skip3",
and
if
found,
delete
that
line
and
the
next
two
lines.

```

#!/bin/sh
sed '/skip3/ {
    N
    N
    s/skip3\n.*\n.*/# 3 lines deleted/
}'

```

Note
that
it
doesn't
matter
what
the
next
two
lines
are.
If
you
wanted
to
match
3
particular
lines,
it's
a
little
more
work.

This
script
looks
for
three
lines,
where
the
first
line
contains
"one",
the
second
contained
"two"
and
the
third
contains
"three",
and
if
found,
replace
them
with
the
string
"1+2+3":

```
#!/bin/sh
sed '
  /one/ {
    N
    /two/ {
      N
      /three/ {
        N
        s/one\ntwo\nthree/1+2+3/
      }
    }
  }
'
```

Matching patterns that span multiple lines

You
can
either
search
for
a
particular
pattern
on
two
consecutive
lines,
or
you
can
search
for
two
consecutive
words
that
may
be
split
on
a
line
boundary.
The
next
example
will
look
for
two
words
which
are
either
on
the
same
line
or
one
is
on
the
end
of
a
line
and
the
second
is
on

the
beginning
of
the
next
line.
If
found,
the
first
word
is
deleted:

```
#!/bin/sh
sed '
/ONE/ {
# append a line
    N
# "ONE TWO" on same line
    s/ONE TWO/TWO/
# "ONE
# TWO" on two consecutive lines
    s/ONE\nTWO/TWO/
}' file
```

Let's
use
the
"D"
command,
and
if
we
find
a
line
containing
"TWO"
immediately
after
a
line
containing
"ONE,"
then
delete
the
first
line:

```
#!/bin/sh
sed '
/ONE/ {
# append a line
    N
# if TWO found, delete the first line
    /\n.*TWO/ D
}' file
```

Click
here
to
get
file:
[sed delete line after word.sh](#)

If
we
wanted

to
print
the
first
line
instead
of
deleting
it,
and
not
print
every
other
line,
change
the
"D"
to
a
"p"
and
add
a
"-n"
as
an
argument
to
sed:

```
#!/bin/sh
sed -n '
# by default - do not print anything
/ONE/ {
# append a line
    N
# if TWO found, print the first line
    /\n.*TWO/ P
}' file
```

Click
here
to
get
file:

[sed_print_line_after_word.sh](#)

It
is
very
common
to
combine
all
three
multi-line
commands.
The
typical
order
is
"N,"
"p"
and
lastly
"D."
This
one

will
delete
everything
between
"ONE"
and
"TWO"
if
they
are
on
one
or
two
consecutive
lines:

```
#!/bin/sh
sed '
/ONE/ {
# append the next line
N
# look for "ONE" followed by "TWO"
/ONE.*TWO/ {
# delete everything between
s/ONE.*TWO/ONE TWO/
# print
P
# then delete the first line
D
}
}' file
```

Click
here
to
get
file:

[sed_delete_between_two_words.sh](#)

Earlier
I
talked
about
the
"="
command,
and
using
it
to
add
line
numbers
to
a
file.
You
can
use
two
invocations
of
sed
to
do
this
(although
it
is

possible
to
do
it
with
one,
but
that
must
wait
until
next
section).
The
first
sed
command
will
output
a
line
number
on
one
line,
and
then
print
the
line
on
the
next
line.
The
second
invocation
of
sed
will
merge
the
two
lines
together:

```
#!/bin/sh
sed '=' file | \
sed '{
    N
    s/\n/ /
}'
```

Click
here
to
get
file:

[sed_merge_two_lines.sh](#)

If
you
find
it
necessary,
you
can
break

one
line
into
two
lines,
edit
them,
and
merge
them
together
again.
As
an
example,
if
you
had
a
file
that
had
a
hexadecimal
number
followed
by
a
word,
and
you
wanted
to
convert
the
first
word
to
all
upper
case,
you
can
use
the
"y"
command,
but
you
must
first
split
the
line
into
two
lines,
change
one
of
the
two,
and
merge
them
together.
That
is,
a


```
line
containing
0x1fff table2

will
be
changed
into
two
lines:

0x1fff
table2

and
the
first
line
will
be
converted
into
upper
case.
I
will
use
tr
to
convert
the
space
into
a
new
line,
and
then
use
sed
to
do
the
rest.
The
command
would
be

./sed_split <file

and
sed_split
would
be:

#!/bin/sh
tr ' ' '\012' |
sed ' {
    y/abcdef/ABCDEF/
    N
    s/\n/ /
}'

Click
here
to
get
file:
sed\_split.sh
```

It
isn't
obvious,
but
sed
could
be
used
instead
of
tr.
You
can
embed
a
new
line
in
a
substitute
command,
but
you
must
escape
it
with
a
backslash.
It
is
unfortunate
that
you
must
use
"*\n*"
in
the
left
side
of
a
substitute
command,
and
an
embedded
new
line
in
the
right
hand
side.
Heavy
sigh.
Here
is
the
example:

```
#!/bin/sh
sed '
s/ /\
/' | \
sed ' {
    y/abcdef/ABCDEF/
    N
    s/\n/ /
```

```
} '
Click
here
to
get
file:
sed\_split\_merge.sh
Sometimes
I
add
a
special
character
as
a
marker,
and
look
for
that
character
in
the
input
stream.
When
found,
it
indicates
the
place
a
blank
used
to
be.
A
backslash
is
a
good
character,
except
it
must
be
escaped
with
a
backslash,
and
makes
the
sed
script
obscure.
Save
it
for
that
guy
who
keeps
asking
dumb
questions.
```

The
sed
script
to
change
a
blank
into
a
"\ "
following
by
a
new
line
would
be:

```
#!/bin/sh
sed
's/
/\\"
/'
file
```

Click
here
to
get
file:

[sed_addslash_before_blank.sh](#)

Yeah.
That's
the
ticket.
Or
use
the
C
shell
and
really
confuse
him!

```
#!/bin/csh
-f
sed
\'
s/
/\\"
/'
file
```

Click
here
to
get
file:

[sed_addslash_before_blank.csh](#)

A
few
more
examples
of

that,
and
he'll
never
ask
you
a
question
again!
I
think
I'm
getting
carried
away.
I'll
summarize
with
a
chart
that
covers
the
features
we've
talked
about:

Pattern Space	Next Input	Command	Output	New Pattern Space	New Text Input
AB	CD	n	<default>	CD	EF
AB	CD	N	-	AB\nCD	EF
AB	CD	d	-	-	EF
AB	CD	D	-	-	EF
AB	CD	p	AB	AB	CD
AB	CD	P	AB	AB	CD
AB\nCD	EF	n	<default>	EF	GH
AB\nCD	EF	N	-	AB\nCD\nEF	GH
AB\nCD	EF	d	-	EF	GH
AB\nCD	EF	D	-	CD	EF
AB\nCD	EF	p	AB\nCD	AB\nCD	EF
AB\nCD	EF	P	AB	AB\nCD	EF

**Using
newlines
in
sed
scripts**

Occasionally
one
wishes
to
use
a
new
line
character
in
a
sed
script.

Well,
this
has
some
subtle
issues
here.
If
one
wants
to
search
for
a
new
line,
one
has
to
use
"
n."
Here
is
an
example
where
you
search
for
a
phrase,
and
delete
the
new
line
character
after
that
phrase
-
joining
two
lines
together.

```
(echo a;echo x;echo y) | sed '/x$/ {  
N  
s:x\n:x:  
'
```

which
generates

```
a  
xy
```

However,
if
you
are
inserting
a
new
line,
don't
use
"
n"
-
instead

```
insert
a
literal
new
line
character:

(echo a;echo x;echo y) | sed 's:x:X\
:'
```

generates

```
a
x
y
```

The Hold Buffer

So far we have talked about three concepts of *sed*:

(1) The input stream or data before it is modified,

(2) the output stream or data after it has been modified, and

(3) the pattern space, or buffer containing characters that can be modified and send to

the
output
stream.

There
is
one
more
"location"
to
be
covered:
the
hold
buffer
or
hold
space.
Think
of
it
as
a
spare
pattern
buffer.
It
can
be
used
to
"copy"
or
"remember"
the
data
in
the
pattern
space
for
later.
There
are
five
commands
that
use
the
hold
buffer.

Exchange with X

The
"x"
command
eXchanges
the
pattern
space
with
the
hold

buffer.
By
itself,
the
command
isn't
useful.
Executing
the
sed
command

`sed 'x'`

as
a
filter
adds
a
blank
line
in
the
front,
and
deletes
the
last
line.
It
looks
like
it
didn't
change
the
input
stream
significantly,
but
the
sed
command
is
modifying
every
line.

The
hold
buffer
starts
out
containing
a
blank
line.
When
the
"x"
command
modifies
the
first
line,
line
1
is
saved

in
the
hold
buffer,
and
the
blank
line
takes
the
place
of
the
first
line.
The
second
"x"
command
exchanges
the
second
line
with
the
hold
buffer,
which
contains
the
first
line.
Each
subsequent
line
is
exchanged
with
the
preceding
line.
The
last
line
is
placed
in
the
hold
buffer,
and
is
not
exchanged
a
second
time,
so
it
remains
in
the
hold
buffer
when
the
program
terminates,
and

never
gets
printed.
This
illustrates
that
care
must
be
taken
when
storing
data
in
the
hold
buffer,
because
it
won't
be
output
unless
you
explicitly
request
it.

Example **of** **Context** **Grep**

One
use
of
the
hold
buffer
is
to
remember
previous
lines.
An
example
of
this
is
a
utility
that
acts
like
grep
as
it
shows
you
the
lines
that
match
a
pattern.
In

addition,
it
shows
you
the
line
before
and
after
the
pattern.
That
is,
if
line
8
contains
the
pattern,
this
utility
would
print
lines
7,
8
and
9.

One
way
to
do
this
is
to
see
if
the
line
has
the
pattern.
If
it
does
not
have
the
pattern,
put
the
current
line
in
the
hold
buffer.
If
it
does,
print
the
line
in
the
hold
buffer,
then

the
current
line,
and
then
the
next
line.
After
each
set,
three
dashes
are
printed.
The
script
checks
for
the
existence
of
an
argument,
and
if
missing,
prints
an
error.
Passing
the
argument
into
the
sed
script
is
done
by
turning
off
the
single
quote
mechanism,
inserting
the
"\$1"
into
the
script,
and
starting
up
the
single
quote
again:

```
#!/bin/sh
# grep3 - prints out three lines around pattern
# if there is only one argument, exit

case $# in
    1) ;;
    *) echo "Usage: $0 pattern";exit;;
```

```

esac;
# I hope the argument doesn't contain a /
# if it does, sed will complain

# use sed -n to disable printing
# unless we ask for it
sed -n '
/$1/' !{
    #no match - put the current line in the hold buffer
    x
    # delete the old one, which is
    # now in the pattern buffer
    d
}
/$1/' {
    # a match - get last line
    x
    # print it
    p
    # get the original line back
    x
    # print it
    p
    # get the next line
    n
    # print it
    p
    # now add three dashes as a marker
    a\
---
    # now put this line into the hold buffer
    x
}

```

Click
here
to
get
file:

grep3.sh

You
could
use
this
to
show
the
three
lines
around
a
keyword,
i.e.:

```
grep3 vt100 </etc/termcap
```

Hold
with
h
or
H

The
"x"
command
exchanges
the
hold
buffer

and
the
pattern
buffer.
Both
are
changed.
The
"h"
command
copies
the
pattern
buffer
into
the
hold
buffer.
The
pattern
buffer
is
unchanged.
An
identical
script
to
the
above
uses
the
hold
commands:

```
#!/bin/sh
# grep3 version b - another version using the hold commands
# if there is only one argument, exit

case $# in
    1) ;;
    *) echo "Usage: $0 pattern"; exit ;;
esac;

# again - I hope the argument doesn't contain a /

# use sed -n to disable printing

sed -n '
'/$1/' !{
    # put the non-matching line in the hold buffer
    h
}
'/$1/' {
    # found a line that matches
    # append it to the hold buffer
    H
    # the hold buffer contains 2 lines
    # get the next line
    n
    # and add it to the hold buffer
    H
    # now print it back to the pattern space
    x
    # and print it.
    p
    # add the three hyphens as a marker
    a\
---
}'
```

Click

here
to
get
file:

grep3a.sh

Keeping
more
than
one
line
in
the
hold
buffer

The
"H"
command
allows
you
to
combine
several
lines
in
the
hold
buffer.
It
acts
like
the
"N"
command
as
lines
are
appended
to
the
buffer,
with
a
"\n"
between
the
lines.
You
can
save
several
lines
in
the
hold
buffer,
and
print
them
only
if

a particular pattern is found later.

As an example, take a file that uses spaces as the first character of a line as a continuation character. The files */etc/termcap*, */etc/printcap*, *makefile* and mail messages use spaces or tabs to indicate a continuing of an entry. If you wanted to print the entry before a word, you could use this script. I use a "[^]I" to indicate an actual

```

tab
character:

#!/bin/sh
# print previous entry
sed -n '
/[ ^I]/{
    # line does not start with a space or tab,
    # does it have the pattern we are interested in?
    '/$1/' {
        # yes it does. print three dashes
        i\
---
        # get hold buffer, save current line
        x
        # now print what was in the hold buffer
        p
        # get the original line back
        x
    }
    # store it in the hold buffer
    h
}
# what about lines that start
# with a space or tab?
/[ ^I]/ {
    # append it to the hold buffer
    H
}'

```

Click
here
to
get
file:

[**grep_previous.sh**](#)

You
can
also
use
the
"H"
to
extend
the
context
grep.
In
this
example,
the
program
prints
out
the
two
lines
before
the
pattern,
instead
of
a
single
line.
The
method
to
limit

this
to
two
lines
is
to
use
the
"s"
command
to
keep
one
new
line,
and
deleting
extra
lines.
I
call
it
grep4:

```
#!/bin/sh

# grep4: prints out 4 lines around pattern
# if there is only one argument, exit

case $# in
    1) ;;
    *) echo "Usage: $0 pattern"; exit ;;
esac;

sed -n '
/$1/' !{
    # does not match - add this line to the hold space
    H
    # bring it back into the pattern space
    x
    # Two lines would look like .*\n.*
    # Three lines look like .*\n.*\n.*
    # Delete extra lines - keep two
    s/^\n\(\.*\n.*\)$/\1/
    # now put the two lines (at most) into
    # the hold buffer again
    x
}
/$1/' {
    # matches - append the current line
    H
    # get the next line
    n
    # append that one also
    H
    # bring it back, but keep the current line in
    # the hold buffer. This is the line after the pattern,
    # and we want to place it in hold in case the next line
    # has the desired pattern
    x
    # print the 4 lines
    p
    # add the mark
    a\
---
}'
```

Click
here
to
get
file:

grep4.sh

You
can
modify
this
to
print
any
number
of
lines
around
a
pattern.
As
you
can
see,
you
must
remember
what
is
in
the
hold
space,
and
what
is
in
the
pattern
space.
There
are
other
ways
to
write
the
same
routine.

**Get
with
g
or
G**

Instead
of
exchanging
the
hold
space
with
the
pattern
space,
you
can
copy
the
hold

space
to
the
pattern
space
with
the
"g"
command.
This
deletes
the
pattern
space.
If
you
want
to
append
to
the
pattern
space,
use
the
"G"
command.
This
adds
a
new
line
to
the
pattern
space,
and
copies
the
hold
space
after
the
new
line.

Here
is
another
version
of
the
"grep3"
command.
It
works
just
like
the
previous
one,
but
is
implemented
differently.
This
illustrates
that
sed

has
more
than
one
way
to
solve
many
problems.
What
is
important
is
you
understand
your
problem,
and
document
your
solution:

```
#!/bin/sh
# grep3 version c: use 'G' instead of H

# if there is only one argument, exit

case $# in
    1);;
    *) echo "Usage: $0 pattern";exit;;
esac;

# again - I hope the argument doesn't contain a /

sed -n '
'/ $1/' !{
    # put the non-matching line in the hold buffer
    h
}
'/ $1/' {
    # found a line that matches
    # add the next line to the pattern space
    N
    # exchange the previous line with the
    # 2 in pattern space
    x
    # now add the two lines back
    G
    # and print it.
    p
    # add the three hyphens as a marker
    a\
---
    # remove first 2 lines
    s/.*\n.*\n\(.*\)$/\1/
    # and place in the hold buffer for next time
    h
}'
```

Click
here
to
get
file:

[grep3c.sh](#)

The
"G"
command
makes

it
easy
to
have
two
copies
of
a
line.
Suppose
you
wanted
to
the
convert
the
first
hexadecimal
number
to
uppercase,
and
don't
want
to
use
the
script
I
described
in
an
earlier
column

```
#!/bin/sh
# change the first hex number to upper case format
# uses sed twice
# used as a filter
# convert2uc <in >out
sed '
s/ /\
/' | \
sed ' {
    y/abcdef/ABCDEF/
    N
    s/\n/ /
}'
```

Click
here
to
get
file:
[convert2uc.sh](#)

Here
is
a
solution
that
does
not
require
two
invocations
of
sed:

```
#!/bin/sh
# convert2uc version b
# change the first hex number to upper case format
# uses sed once
# used as a filter
# convert2uc <in >out
sed '
{
    # remember the line
    h
    #change the current line to upper case
    y/abcdef/ABCDEF/
    # add the old line back
    G
    # Keep the first word of the first line,
    # and second word of the second line
    # with one humongous regular expression
    s/^\([^ ]*\) .* \n[^\ ]* \(.*/\1 \2/
}'
```

Click
here
to
get
file:

[**convert2uc1.sh**](#)

Carl
Henrik
Lunde
suggested
a
way
to
make
this
simpler.
I
was
working
too
hard.

```
#!/bin/sh
# convert2uc version b
# change the first hex number to upper case format
# uses sed once
# used as a filter
# convert2uc <in >out
sed '
{
    # remember the line
    h
    #change the current line to upper case
    y/abcdef/ABCDEF/
    # add the old line back
    G
    # Keep the first word of the first line,
    # and second word of the second line
    # with one humongous regular expression
    s/ .* / / # delete all but the first and last word
}'
```

Click
here
to
get
file:

[**convert2uc2.sh**](#)

This
example
only
converts

the
letters
"a"
through
"f"
to
upper
case.
This
was
chosen
to
make
the
script
easier
to
print
in
these
narrow
columns.
You
can
easily
modify
the
script
to
convert
all
letters
to
uppercase,
or
to
change
the
first
letter,
second
word,
etc.

Flow **Control**

As
you
learn
about
sed
you
realize
that
it
has
its
own
programming
language.
It
is
true
that
it's
a

very
specialized
and
simple
language.
What
language
would
be
complete
without
a
method
of
changing
the
flow
control?
There
are
three
commands
sed
uses
for
this.
You
can
specify
a
label
with
an
text
string
preceded
by
a
colon.
The
"b"
command
branches
to
the
label.
The
label
follows
the
command.
If
no
label
is
there,
branch
to
the
end
of
the
script.
The
"t"
command
is
used
to

test
conditions.
Before
I
discuss
the
"t"
command,
I
will
show
you
an
example
using
the
"b"
command.

This
example
remembers
paragraphs,
and
if
it
contains
the
pattern
(specified
by
an
argument),
the
script
prints
out
the
entire
paragraph.

```
#!/bin/sh
sed -n '
# if an empty line, check the paragraph
/^$/ b para
# else add it to the hold buffer
H
# at end of file, check paragraph
$ b para
# now branch to end of script
b
# this is where a paragraph is checked for the pattern
:para
# return the entire paragraph
# into the pattern space
x
# look for the pattern, if there - print
/ '$1' / p
'
```

Click
here
to
get
file:

[**grep_paragraph.sh**](#)

Testing
with

t

You can execute a branch if a pattern is found. You may want to execute a branch only if a substitution is made. The command "t label" will branch to the label if the last substitute command modified the pattern space.

One use for this is recursive patterns. Suppose you wanted to remove white space inside parenthesis. These parentheses might be nested. That is, you

```
would
want
to
delete
a
string
that
looked
like
"(
(
(
)))
)."
The
sed
expressions

sed 's/([ ^I]*)/g'

would
only
remove
the
innermost
set.
You
would
have
to
pipe
the
data
through
the
script
four
times
to
remove
each
set
or
parenthesis.
You
could
use
the
regular
expression

sed 's/([ ^I()*)/g'

but
that
would
delete
non-matching
sets
of
parenthesis.
The
"t"
command
would
solve
this:

#!/bin/sh
sed '
```

```
:again  
s/([ ^I]*)//  
t again  
,
```

An
earlier
version
had
a
'g'
after
the
's'
expression.
This
is
not
needed.

Click
here
to
get
file:
[delete_nested_parens.sh](#)

Debugging **with** **!**

The
'!'
command
will
print
the
pattern
space
in
an
unambiguous
form.
Non-printing
characters
are
printed
in
a
C-style
escaped
format.

This
can
be
useful
when
debugging
a
complex
multi-line
sed
script.

An

alternate way of adding comments

There
is
one
way
to
add
comments
in
a
sed
script
if
you
don't
have
a
version
that
supports
it.
Use
the
"a"
command
with
the
line
number
of
zero:

```
#!/bin/sh
sed '
/begin/ {
0i\
    This is a comment\
    It can cover several lines\
    It will work with any version of sed
}'
```

Click
here
to
get
file:

[sed_add_comments.sh](#)

The poorly documented

i

There
is
one
more

```
sed
command
that
isn't
well
documented.
It
is
the
";"
command.
This
can
be
used
to
combined
several
sed
commands
on
one
line.
Here
is
the
grep4
script
I
described
earlier,
but
without
the
comments
or
error
checking
and
with
semicolons
between
commands:

#!/bin/sh
sed
-n
'
/$1/'
!{;H;x;s/^.*\n
\(. *\n.*
\)$/\1/;x;}
/$1/'
{;H;n;H;x;p;a\
---
}'

Click
here
to
get
file:
grep4a.sh

Yessireebob!
Definitely
character
building.
```


I
think
I
have
made
my
point.
As
far
as
I
am
concerned,
the
only
time
the
semicolon
is
useful
is
when
you
want
to
type
the
sed
script
on
the
command
line.
If
you
are
going
to
place
it
in
a
script,
format
it
so
it
is
readable.
I
have
mentioned
earlier
that
many
versions
of
sed
do
not
support
comments
except
on
the
first
line.
You
may

want
to
write
your
scripts
with
comments
in
them,
and
install
them
in
"binary"
form
without
comments.
This
should
not
be
difficult.
After
all,
you
have
become
a
sed
guru
by
now.
I
won't
even
tell
you
how
to
write
a
script
to
strip
out
comments.
That
would
be
insulting
your
intelligence.
Also
-
some
operating
systems
do
NOT
let
you
use
semicolons.
So
if
you
see
a
script

with
semicolons,
and
it
does
not
work
on
a
non-Linux
system,
replace
the
semicolon
with
a
new
line
character.
(As
long
as
you
are
not
using
csh/tcsh,
but
that's
another
topic.

Passing regular expressions as arguments

In
the
earlier
scripts,
I
mentioned
that
you
would
have
problems
if
you
passed
an
argument
to
the
script
that
had
a
slash
in
it.
In
fact,

```
regular
expression
might
cause
you
problems.
A
script
like
the
following
is
asking
to
be
broken
some
day:

#!/bin/sh
sed 's/""$1""//g'

If
the
argument
contains
any
of
these
characters
in
it,
you
may
get
a
broken
script:
"/\.*[]^$"
For
instance,
if
someone
types
a
"/"
then
the
substitute
command
will
see
four
delimiters
instead
of
three.
You
will
also
get
syntax
errors
if
you
provide
a
"]"
without
```

```
a
"]".
One
solution
is
to
have
the
user
put
a
backslash
before
any
of
these
characters
when
they
pass
it
as
an
argument.
However,
the
user
has
to
know
which
characters
are
special.
Another
solution
is
to
add
a
backslash
before
each
of
those
characters
in
the
script

#!/bin/sh
arg=`echo "$1" | sed 's:[ ]\[^\$\\.\\*\\/]:\\\\&:g'`
sed 's/\'$arg\'//g'

Click
here
to
get
file:
sed\_with\_regular\_expressions1.sh
If
you
were
searching
for
the
pattern
"^./,"
```

the
script
would
convert
this
into
"^\.\.V"
before
passing
it
to
sed.

Inserting binary characters

Dealing
with
binary
characters
can
be
trick,
expecially
when
writing
scripts
for
people
to
read.
I
can
insert
a
binary
character
using
an
editor
like
EMACS
but
if
I
show
the
binary
character,
the
terminal
may
change
it
to
show
it
to
you.

The
easiest
way
I
have

```
found
to
do
this
in
a
script
in
a
portable
fashion
is
to
use
the
tr(1)
command.
It
understands
octal
notations,
and
it
can
be
output
into
a
variable
which
can
be
used.

Here's
a
script
that
will
replace
the
string
"ding"
with
the
ASCII
bell
character:

#!/bin/sh
BELL=`echo x | tr 'x' '\007'`
sed "s/ding/$BELL/"

Please
note
that
I
used
double
quotes.
Since
special
characters
are
interpreted,
you
have
to
be
careful
```

when
you
use
this
mechanism.

GNU sed Command Line arguments

One
of
the
conventions
UNIX
systems
have
is
to
use
single
letters
are
command
line
arguments.
This
makes
typing
faster,
and
shorted,
which
is
an
advantage
if
you
are
in
a
contest.
Normal
people
often
find
sed's
terseness
cryptic.
You
can
improve
the
readability
of
sed
scripts
by
using
the
long
word
equivalent

options.
That
is,
instead
of
typing

```
sed -n 20p
```

You
can
type
the
long
word
version
of
the
-n
argument

```
sed --quiet 20p
```

Or

```
sed --silent 20p
```

The
long
form
of
sed's
command
line
arguments
always
have
2
hyphens
before
their
names.
GNU
sed
has
the
following
long-form
command
line
arguments:

GNU Command Line Arguments	
Short Form	Long Form
-n	--quiet --silent
-e script	--expression=SCRIPT
-f SCRIPTFILE	--file=SCRIPTFILE

-i[SUFFIX]	--in-place[=SUFFIX]
-l N	--line-length=N
	--posix
-b	--binary
	--follow-symlinks
-r	--regular-extended
-s	--separate
-u	--unbuffered
	--help
	--version

Let's
define
each
of
these.

The
-posix
argument

The
GNU
version
of
sed
has
many
features
that
are
not
available
in
other
versions.
When
portability
is
important,
test
your
script
with
the

-posix
option.
If
you
had
an
example
that
used
a
feature
of
GNU
sed,
such
as
the
'v'
command
to
test
the
version
number,
such
as

```
#this is a sed command file  
v 4.0.1  
# print the number of lines  
$=
```

And
you
executed
it
with
the
command

```
sed -nf sedfile --posix <file
```

then
the
GNU
version
of
sed
program
would
give
you
a
warning
that
your
sed
script
is
not
compatible.
It
would
report:

```
sed: -e expression #1, char 2: unknown command: `v'
```

The --version argument

You
can

determine
which
version
of
sed
you
are
using
with
the
GNU
sed
--version
command.
This
is
what
it
outputs
on
my
computer

```
# sed --version
GNU sed version 4.2.1
Copyright (C) 2009 Free Software Foundation, Inc.
This is free software; see the source for copying conditions.  There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE,
to the extent permitted by law.

GNU sed home page: <http://www.gnu.org/software/sed/>.
General help using GNU software: <http://www.gnu.org/gethelp/>.
E-mail bug reports to: <bug-gnu-utils@gnu.org>.
Be sure to include the word ``sed'' somewhere in the ``Subject:'' field.
```

The **-h** **Help** **argument**

The
-h
option
will
print
a
summary
of
the
sed
commands.
The
long
argument
of
the
command
is
sed --help

It
provides
a
nice
summary
of
the
command
line

arguments.

The -l Line Length Argument

I've already described the 'l' command. The default line width for the 'l' command is 70 characters. This default value can be changed by adding the '-l' N' option and specifying the maximum line length as the number after the '-l'.

```
sed -n -l 80 'l' <file
```

The long form version of the command line is

```
sed -n --line-length=80 'l' <file
```

The -s

Separate argument

Normally, when you specify several files on the command line, sed concatenates the files into one stream, and then operates on that single stream.

If you had three files, each with 100 lines, then the command

```
sed -n '1,10 p' file1 file2 file3
```

would only print the first 10 lines of file1. The '-s' command tells GNU sed to treat the files as independent files, and to print out the first 10 lines

of
each
file,
which
is
similar
to
the
head
command.

Here's
another
example:

If
you
wanted
to
print
the
number
of
lines

of
each
file,
you
could
use

`'wc
-l'`

which
prints
the
number
of

lines,
and
the
filename,
for

each
file,
and
at

the
end
print
the

total
number
of

lines.

Here

is

a

simple

shell

script

that

does

something

similar,

just

using

sed:

```
#!/bin/sh
```

```
FILES=$*
```

```
sed -s -n '$=' $FILES # print the number of lines for each file
```

```
sed -n '$=' $FILES # print the total number of lines.
```

The

`'wc`

`-l'`

command

does

print

out

the

filenames,

unlike

the

above

```

script.
A
better
emulation
of
the
'wc
-i'
command
would
execute
the
command
in
a
loop,
and
print
the
filenames.
Here
is
a
more
advanced
script
that
does
this,
but
it
doesn't
use
the
'-s'
command:

#!/bin/sh
for F in "$@"
do
    NL=`sed -n '$=' < "$F" ` && printf "  %d %s\n" $NL "$F"
done
TOTAL=`sed -n '$=' "$@"`
printf "  %d total\n" $TOTAL

```

The -i in-place argument

I've
 already
 described
 in
 Editing
 multiple
 files
 the
 way
 I
 like
 to
 do
 this.
 For
 those
 who
 want
 a
 simpler
 method,
 GNU
 Sed


```
allows
you
to
do
this
with
a
command
line
option
-
"-i".
Let's
assume
that
we
are
going
to
make
the
same
simple
change
-
adding
a
tab
before
each
line.
This
is
a
way
to
do
this
for
all
files
in
a
directory
with
the
".txt"
extension
in
the
current
directory:

sed -i 's/^\t/' *.txt

The
long
argument
name
version
is

sed --in-place 's/^\t/' *.txt

This
version
deletes
the
original
file.
If
you
```

```
are
as
cautious
as
I
am,
you
may
prefer
to
specify
an
extension,
which
is
used
to
keep
a
copy
of
the
original:

sed -i.tmp 's/^/t/' *.txt

And
the
long
argument
name
version
is

sed --in-place=.tmp 's/^/t/' *.txt

In
the
last
two
versions,
the
original
version
of
the
"a.txt"
file
would
have
the
name
"a.txt.tmp".
You
can
then
delete
the
original
files
after
you
make
sure
all
worked
as
you
expected.
Please
consider
the
backup
option,
and
heed
my
warning.
You
can
easily
delete
the
```

backed-up
original
file,
as
long
as
the
extension
is
unique.

The
GNU
version
of
sed
allows
you
to
use
"-i"
without
an
argument.
The
FreeBSD/Mac
OS
X
does
not.
You

must
provide
an
extension
for
the
FreeBSD/Mac
OS
X
version.
If
you
want
to
do
in-place
editing
without
creating
a
backup,
you
can
use

```
sed -i '' 's/^\t/' *.txt
```

The --follow- symlinks argument

The
in-place
editing
feature
is

handy
to
have.
But
what
happens
if
the
file
you
are
editing
is
a
symbolic
link
to
another
file?
Let's
assume
you
have
a
file
named
"b"
in
a
directory
called
"tmp",
with
a
symbolic
link
to
this
file:

```
$ ls -l b  
lrwxrwxrwx 1 barnett adm 6 Mar 16 16:03 b.txt -> tmp/b.txt
```

If
you
executed
the
above
command
to
do
in
place
editing,
there
will
be
a
new
file
called
"b.txt"
in
the
current
directory,
and
"tmp/b.txt"
will
be
unchanged.
Now
you
have
two

```
versions
of
the
file,
one
is
changed
(in
the
current
directory),
and
one
is
not
(in
the
"tmp"
directory).
And
where
you
had
a
symbolic
link,
it
has
been
replaced
with
a
modified
version
of
the
original
file.
If
you
want
to
edit
the
real
file,
and
keep
the
symbolic
link
in
place,
use
the
"--follow-
symlinks"
command
line
option:

sed -i --follow-symlinks 's/^\t/' *.txt

This
follows
the
symlink
to
the
original
location,
and
modifies
the
file
in
the
"tmp"
directory,
If
you
specify
an
```

extension,
the
original
file
will
be
found
with
that
extension
in
the
same
directory
as
the
real
source.
Without
the
--follow-
symlinks
command
line
option,
the
"backup"
file
"b.tmp"
will
be
in
the
same
directory
that
held
the
symbolic
link,
and
will
still
be
a
symbolic
link
-
just
renamed
to
give
it
a
new
extension.

The **-b** **Binary** **argument**

Unix
and
Linux
systems
consider
the
new
line
character
"\n"
to
be
the

```
end
of
the
line.
However,
MS-DOS,
Windows,
and
Cygwin
systems
end
each
line
with
"\r\n"
-
Carriage
return
and
line-feed.
If
you
are
using
any
of
these
operating
systems,
the
"-b"
or
"--binary"
command
line
option
will
treat
the
carriage
return/new
line
combination
as
the
end
of
the
line.
Otherwise
the
carriage
return
is
treated
as
an
unprintable
character
immediately
before
the
end-of-line.
I
think.
(Note
to
self
-
```

verify
this).

The -r Extended Regular Expression argument

When
I
mention
patterns,
such
as
"s/pattern/",
the
pattern
is
a
regular
expression.
There
are
two
common
classes
of
regular
expressions,
the
original
"basic"
expressions,
and
the
"extended"
regular
expressions.

For
more
on
the
differences
see

My
tutorial
on
regular
expressions

and
the
the
section
on
extended
regular
expressions .

Because
the
meaning
of
certain

characters
are
different
between
the
regular
and
extended
expressions,
you
need
a
command
line
argument
to
enable
sed
to
use
the
extension.
To
enable
this
extension,
use
the
"-r"
command,
as
mentioned
in

the
example
on
finding
duplicated
words
on
a
line

```
sed -r -n '/\([a-z]+\) \1/p'
```

or

```
sed --regular-extended -quiet '/\([a-z]+\) \1/p'
```

I
already
mentioned
that
Mac
OS
X
and
FreeBSD
uses
-E
instead
of
-r .

The
-u
Unbuffered

argument

Normally
-
Unix
and
Linux
systems
apply
some
intelligence
to
handling
standard
output.
It's
assumed
that
if
you
are
sending
results
to
a
terminal,
you
want
the
output
as
soon
as
it
becomes
available.
However,
if
you
are
sending
the
output
to
a
file,
then
it's
assumed
you
want
better
performance,
so
it
buffers
the
output
until
the
buffer
is
full,
and
then
the
contents
of

the
buffer
is
written
to
the
file.
Let
me
elaborate
on
this.
Let's
assume
for
this
example
you
have
a
very
large
file,
and
you
are
using
sed
to
search
for
a
string,
and
to
print
it
when
it
is
found:

```
sed -n '/MATCH/p' <file
```

Since
the
output
is
the
terminal,
as
soon
as
a
match
is
found,
it
is
printed.
However,
if
sed
pipes
its
output
to
another
program,
it
will
buffer
the
results.

But
there
are
times
when
you
want
immediate
results.
This
is
especially
true
when
you
are
dealing
with
large
files,
or
files
that
occasionally
generate
data.
To
summarize,
you
have
lots
of
input
data,
and
you
want
sed
to
process
it,
and
then
send
this
to
another
program
that
processes
the
results,
but
you
want
the
results
when
it
happens,
and
not
delayed.
Let
me
make
up
a
simple
example.
It's
contrived,
but
it
does
explain
how
this
works.
Here's
a
program

```
called
SlowText
that
prints
numbers
from
1
to
60,
once
a
second:

#!/bin/sh
for i in `seq 1 60`
do
    echo $i
    sleep 1
done
```

Let's
use
sed
to
search
for
lines
that
have
the
character
'1',
and
have
it
send
results
to
awk,
which
will
calculate
the
square
of
that
number.
This
would
be
the
admittedly
contrived
script:

```
SlowText | sed -n '/1/p' | awk '{print $1*$1}'
```

This
works,
but
because
sed
is
buffering
the
results,
we
have
to
wait
until
the
buffer
fills
up,
or
until
the
SlowText
program
exists,
before
we

the
results.
You
can
eliminate
the
buffering,
and
see
the
results
as
soon
as
SlowText
outputs
them,
by
using
the
"-u"
option.
With
this
option,
you
will
see
the
squares
printed
as
soon
as
possible:

```
SlowText | sed -un '/1/p' | awk '{print $1*$1}'
```

The
long
form
of
the
argument
is
"--unbuffered".

Mac
OS
X
and
FreeBSD
use
the
argument
"-l".

GNU
Sed
4.2.2
and
later
will
also
be
unbuffered
while
reading
files,
not
just
writing
them.

The **-Z** **Null** **Data** **argument**

Normally,
sed
reads
a
line
by
reading
a
string
of
characters
up
to
the
end-of-line
character
(new
line
or
carriage
return).
See

the **-b** **Binary** **command** **line** **argument**

The
GNU
version
of
sed
added
a
feature
in
version
4.2.2
to
use
the
"NULL"
character
instead.
This
can
be
useful
if
you
have
files
that
use
the
NULL
as
a

record
separator.
Some
GNU
utilities
can
generate
output
that
uses
a
NULL
instead
a
new
line,
such
as
"find
.
-print0"
or
"grep
-lZ".
This
feature
is
useful
if
you
are
operating
on
filenames
that
might
contain
spaces
or
binary
characters.

For
instance,
if
you
wanted
to
use
"find"
to
search
for
files
and
you
used
the
"-print0"
option
to
print
a
NULL
at
the
end
of
each

filename,
you
could
use
sed
to
delete
the
directory
pathname:

```
find . -type f -print0 | sed -z 's:^.*/:::' | xargs -0 echo
```

The
above
example
is
not
terribly
useful
as
the
"xargs"
use
of
echo
does
not
retain
the
ability
to
retain
spaces
as
part
of
the
filename.
But
is
does
show
how
to
use
the
sed
"-z"
command.

GNU
grep
also
has
a
-Z
option
to
search
for
strings
in
files,
placing
a
"NULL"
at
the

end
of
each
filename
instead
of
a
new
line.
And
with
the
-l
command,
grep
will
print
the
filename
that
contains
the
string,
retaining
non-printing
and
binary
characters:

```
grep -lZ STRING */** | sed -z 's:^.*/::' | xargs -0 echo
```

This
feature
is
very
useful
when
users
have
the
ability
to
create
their
own
filenames.

FreeBSD **Extensions**

Apple
uses
the
FreeBSD
version
of
sed
for
Mac
OS
X
instead
of
the
GNU
sed.
However,

the
FreeBSD
version
has
a
couple
of
additions.

The -a or delayed open Argument

Normally,
as
soon
as
sed
starts
up,
it
opens
all
files
that
are
referred
to
by
the
"w"
command.
The
FreeBSD
version
of
sed
has
an
option
to
delay
this
action
until
the
"w"
command
is
executed.

The -I in-place argument

FreeBSD
added

a
"-I"
option
that
is
similar
to
the
-i
option.
The
"-i"
option
treats
the
editing
each
file
as
a
separate
instance
of
sed.
If
the
"-I"
option
is
used,
then
line
numbers
do
not
get
reset
at
the
beginning
of
each
line,
and
ranges
of
addresses
continue
from
one
file
to
the
next.
That
is,
if
you
used
the
range
'/BEGIN
/,/END/'
and
you
used
the
"-I"
option,

you
can
have
the
"BEGIN"
in
the
first
file,
and
"END"
in
the
second
file,
and
the
commands
executed
within
the
range
would
span
both
files.
If
you
used
"-i",
then
the
commands
would
not.

And
like
the
-i
option,
the
extension
used
to
store
the
backup
file
must
be
specified.

-E **or** **Extended** **Regular** **Expressions**

I
mentioned
extended
regular
expressions
earlier .

FreeBSD
(and
Mac
OS
X)
uses
"-E"
to
enable
this.
However,
FreeBSD
later
added
the
-r
command
to
be
compatible
with
GNU
sed.

Using word boundaries

Someone
once
asked
me
to
help
them
solve
a
tricky
sed
problem
involving
word
boundaries.
Let's
suppose
you
have
the
following
input

```
/usr/bin /usr/local/bin /usr/local /usr/local/project/bin
```

and
you
wanted
to
delete
"/usr/local"
but
leave
the
other
3
paths
alone.
You

could
use
the
simple
(and
incorrect)
command:

```
sed 's@/usr/local@@'
```

which
would
output

```
/usr/bin /bin /usr/local /usr/local/project/bin
```

That
is,
it
would
mistakenly
change
'/usr/local/bin'
to
'/bin'
and
not
delete
'/usr/local'
which
was
the
intention
of
the
programmer.
The
better
method
is
to
include
spaces
around
the
search:

```
sed 's@ /usr/local @ @'
```

However,
this
won't
work
if
'/usr/local'
is
at
the
beginning,
or
at
the
end
of
the
line.
It
also
won't
work

```
if
'/usr/local'
is
the
only
path
on
the
line.
To
handle
these
edge
cases,
you
can
simply
describe
all
of
these
conditions
as
separate
cases:

#!/bin/sh
sed '
s@ /usr/local @ @g
s@^/usr/local @@
s@ /usr/local$@@
s@^/usr/local$@@
'
```

This works fine if the string you are searching for is surrounded by a space. But what happens if the string is surrounded by other characters, which may be one of several possible characters? You can


```
always
make
up
your
own
class
of
characters
that
define
the
'end
of
a
word';
For
instance,
if
your
string
consists
of
alphanumeric
characters
and
the
slash,
the
class
of
characters
can
be
defined
by
'[a-zA-
Z0-9/]'
or
the
more
flexible
'[[:alnum:]]/'.
We
can
define
the
class
fo
characters
to
be
all
but
these,
by
using
the
caret,
i.e.
'^[[:alnum:]]/'.
And
unlike
the
space
before,
if
you
are
going
```

to
 use
 character
 classes,
 you
 may
 have
 to
 remember
 what
 these
 characters
 are
 and
 not
 delete
 them.
 So
 we
 can
 replace
 the
 space
 with
 '[^[:alnum:]]/'
 and
 then
 change
 the
 command
 to
 be

```
#!/bin/sh
sed '
s@\[^[[:alnum:]]\]/usr/local\[^[[:alnum:]]\]/@l12@g
s^/usr/local\[^[[:alnum:]]\]/@l1@
s@\[^[[:alnum:]]\]/usr/local$@l1@
s^/usr/local$@@
'
```

The
 first
 version
 would
 replace

/usr/local

with
 a
 single
 space.

This
 method
 would
 replace
 '/usr
 /local:'
 with
 ':'

because
 the
 redundant
 deliniators
 are
 not
 deleted.
 Be

sure
to
fix
this
if
you
need
to.

This
method
always
works,
but
it
is
inelegant
and
error
prone.
There
are
other
methods,
but
they
**may
not
be
portable.**

Solaris's
version
of
sed
used
the
special
characters
'\<>'
and
'\<>'
as
anchors
that
indicated
a
word
boundary.
So
you
could
use

```
s@\<>/usr/local\<>@@
```

However,
the
GNU
version
of
sed
says
the
usage
of
these
special
characters
are

undefined.
According
to
the
manual
page:

```
Regex syntax clashes (problems with backslashes)
`sed' uses the POSIX basic regular expression syntax.  According to
the standard, the meaning of some escape sequences is undefined in
this syntax; notable in the case of `sed' are `\\|', `\\+', `\\?',
`\\', `\\'', `\\<', `\\>', `\\b', `\\B', `\\w', and `\\W'.

As in all GNU programs that use POSIX basic regular expressions,
`sed' interprets these escape sequences as special characters.
So, `x\\+' matches one or more occurrences of `x'.  `abc\\|def'
matches either `abc' or `def'.
```

When
in
doubt,
experiment.

Command Summary

As
I
promised
earlier,
here
is
a
table
that
summarizes
the
different
commands.
The
second
column
specifies
if
the
command
can
have
a
range
or
pair
of
addresses
or
a
single
address
or
pattern.
The
next
four
columns
specifies
which
of
the
four

buffers
 or
 streams
 are
 modified
 by
 the
 command.
 Some
 commands
 only
 affect
 the
 output
 stream,
 others
 only
 affect
 the
 hold
 buffer.
 If
 you
 remember
 that
 the
 pattern
 space
 is
 output
 (unless
 a
 "-n"
 was
 given
 to
sed),
 this
 table
 should
 help
 you
 keep
 track
 of
 the
 various
 commands.

Command	Address or Range	Modification to Input Stream	Modification to Output Stream	Modification to Pattern Space	Modification to Hold Buffer
=	-	-	Y	-	-
a	Address	-	Y	-	-
b	Range	-	-	-	-
c	Range	-	Y	-	-
d	Range	Y	-	Y	-
D	Range	Y	-	Y	-
g	Range	-	-	Y	-
G	Range	-	-	Y	-
h	Range	-	-	-	Y
H	Range	-	-	-	Y
i	Address	-	Y	-	-
I	Address	-	Y	-	-
n	Range	Y	*	-	-
N	Range	Y	-	Y	-

p	Range	-	Y	-	-
P	Range	-	Y	-	-
q	Address	-	-	-	-
r	Address	-	Y	-	-
s	Range	-	-	Y	-
t	Range	-	-	-	-
w	Range	-	Y	-	-
x	Range	-	-	Y	Y
y	Range	-	-	Y	-

The "n" command may or may not generate output, depending on the "-n" option. The "r" command can only have one address, despite the documentation.

Check out my new [Sed Reference Chart](#)

In
Conclusion

This concludes my tutorial on *sed*. It is possible to find shorter forms of some of my scripts. However,

I
chose
these
examples
to
illustrate
some
basic
constructs.
I
wanted
clarity,
not
obscurity.
I
hope
you
enjoyed
it.

More References

This
concludes
my
tutorial
on
sed.
Other
of
my
UNIX
shell
tutorials
can
be
found
here.
Other
shell
tutorials
and
references
can
be
found
at

**FreeBSD
Sed
Man
Page
Apple/Mac
OS
X
Sed
Man
Page
GNU
Sed
Manual
GNU
Sed
4.2.2
Release**

Notes
sed(1)
Seventh
Eddition
Unix
sed(1)
manual
page
from
Sun/Oracle
Heiner's
SHELLdorado
Chris
F.
A.
Johnson's
UNIX
Shell
Page
The
Wikipedia
Entry
on
SED
SED
one-liners

And
don't
forget
The
SED
FAQ

This
document
was
originally
converted
from
NROFF
to
TEXT
to
HTML.
Please
forgive
errors
in
the
translation.
If
you
are
confused,
grab
the
actual
script
if
possible.
No
translations
occurred
in
the
scripts.

**Thanks
for
the
feedback,
gang**

Thanks
to
Keelan
Evans,
Fredrik
Nilsson,
and
Kurt
McKee
for
spotting
some
typos.
Thanks
to
Wim
Stolker
and
Jose'
Sebrosa
as
well.
Thanks
to
Olivier
Mengue.
Thanks
to
Andrew
M.
Goth.
Thanks
to
David
P.
Brown.
Thanks
to
Axel
Schulze
for
some
corrections
Thanks
to
Martin
Jan
for
the
corrections
in
sed
format
(grin)
Thanks
to
David
Ward
for
some

corrections
A
big
thanks
for
Fazl
Rahman
for
spotting
dozens
of
errors.
Thanks
to
Carl
Henrik
Lunde
who
suggested
an
improvement
to
convert2uc1.sh
A
big
thanks
to
Bryan
Hyun
Huh
who
spotted
an
error
in
the
table
and
reference
chart
Thanks
for
input
from

Marten Jan
Gordon Wilson
Tom Konantz
Peter Bratton
Grant Root
Keith Briggs
Zoltan Miklos
Peggy Russell
Lorens Kockkum.net
John Poulin
Rihards
Corey Richardson
Eric Mathison
Ildar Mulyukov
Tom Zhu
Abhijeet Rastogi [@shadyabhi](#)
Steve LeBlanc [@slevo](#)
dontforget yourtowel [@whatissixbynine](#)
Yiming
Fei Wang
Kenneth R. Beesley
Duncan Sung W. Kim [@DuncanSungWKim](#)
Juan Eugenio Abadie
Zander Hill [@ZPH](#)
[Rob Smith](#)
[Peter Moore](#)

*This
document
was
translated
by
troff2html
v0.21
on
September
22,
2001
and
then
manually
edited
to
make
it
compliant
with:*

